

Part 3 Open Addressing Testing Strategy

Paul Nickerson

In addition to separating tests by operations that are expected to fail and those that are expected to succeed, I also included a scenario for testing the `cluster_distribution()` function.

operation_failures_(probe type/double hash).cpp

I include three failure operations scenarios, one for each of the secondary hashing methods to be supported (linear probing, quadratic probing, and double hashing). Within each scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string).

Linear Probing

To the linear probing instance, I start with an empty map and try to `search()` and `remove()` an item whose key does not exist in the map. Both calls should return a value less than zero (indicating key not found). I then fill the map to capacity, and attempt to insert an item where there is no space for it, which is expected to return a value less than zero. From this newly-filled map, I attempt to `remove()` a key that doesn't exist in the map, and `search()` for a key that does not exist in the map, both of which should return a value less than zero.

Quadratic Probing and Double Hashing

I was not able to get the `remove()` methods working for the quadratic probing or the double hashing instances. Quadratic probing tends to have a nasty habit of not reliably visiting every potential slot. Double hashing uses the key itself to pick the next slot, which leads to the following scenario:

- Keys A and B give equivalent primary hashing values but different double hashing values
- Key A exists in slot 1
- Key B is inserted, skipping slot 1 and probing forward to slot 2
- Key A is removed, and afterwards we attempt to resolve the cluster and reposition Key B so that it can be visited. However, Key A's probing value is dependent on Key A, so it is impossible to know that Key B originally intended to take Key A's slot, so

when subsequently searching for Key B we will visit an empty slot where A was and encounter an empty slot

Because of these issues, I did not test the `remove()` functionality of the quadratic probing and double hashing instances, opting instead to simply fill the map to capacity and then search for a key that should not exist in the map. Note that quadratic hashing, as previously mentioned, has a tendency to sometimes ignore some slots, and the probability of this occurring increases drastically as the load factor approaches 1. I early-abandon when this becomes evident to avoid an infinite loop, but, as a result, it was not possible to include a test to try and insert into a completely filled map like I could do in linear probing, since one key may induce an infinite loop while another key successfully finds one of the last remaining slots.

operation__successes__(probe type/double hash).cpp

Within the three success operations scenarios (one for each of the hashing methods), I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the `print()` function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. The number of hyphens in the `print()` output should equal `capacity() - size()`, ie the number of unoccupied slots.

As previously mentioned, I was unable to successfully implement `remove()` functionality in quadratic probing and double hashing, but within linear probing I attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

cluster__distribution__(probe type/double hash).cpp

Since `cluster_distribution()` returns a priority queue of clusters, and each cluster has a minimum size of one, all the clusters taken together should encompass every occupied slot. Therefore, I fill the map with a bunch of items, then clear it and fill it again to try and destabilize the map. From there, I take the summation, over every cluster, of the cluster's size times the number of clusters having that size. The result of that summation should equal the output from the map's `size()` method. I do this four times for each of the hashing method instances, once for each of the key types supported (a total of 12 times).