

Part I

Part 1 was fairly straightforward. Linear probing lets us make several assumptions about its behavior. For one, from any arbitrary array slot, we can probe to every other slot in exactly M iterations (where M is the map capacity). No infinite loops to worry about! The trickiest part of part 1 was ensuring that, once we remove an item from the array, we probe forward and remove/re-insert each item in the cluster. If I needed to implement a quick and dirty hashing implementation, I would likely choose open addressing with linear probing as my first choice unless there was a good reason that it did not work well for my use-case.

Part II

I really enjoyed part 2, simply because I love working with linked lists. At each slot, I hold on to a head pointer to the first item in the chain, and include a tail dummy node at the end, but I don't store a pointer to the tail. This makes it easy to update the chain. For instance, when I insert a new item, I traverse until I either find the key (in which case I simply update that node's value), or encounter the tail dummy node. If I hit the tail, that means the key was not yet stored in the list, and I transform the tail into a regular node holding the new item and append a new tail.

When removing, I simply traverse the list until I either hit the tail (item not found), or find an item. In that case, I replace the current item with its next node using a copy assignment operator, and delete the next node's pointer.

Part III Bucket

Part 3 introduced new key types, which were all straightforward to implement except for the c-strings. I hate passing pointers around, especially if they come from the user and refer to character arrays that may go out of scope and get. What I did to solve this was wrap all keys and values in their own class, Key and Value respectively. These classes use another wrapper class, GenericContainer, to hold their raw value. I use template specialization to treat c-strings as special. When GenericContainer is initialized with a c-string, it makes a copy for itself. Because Keys and Values can be replaced during normal map operations, we don't want to lose the pointers. Therefore, whenever a raw value must persist (ie, when returning the value to the client), I create another copy of the raw value and pass the pointer to the client. The client is then responsible for deleting the c-string if he wishing to avoid memory leaks. This strategy allows me to abstract away the pointer logic, greatly reducing code complexity.

In addition, part III requested a new method, `cluster_distribution`. To implement this, I create an integer array of size M (where M is the map capacity()). Each index in that array refers to a cluster size, and the associated value is used as a counter to track the number of encountered cluster instances of that size. Once I take inventory of every cluster, I traverse that array and pass extant clusters to a priority queue and return that, which sorts the clusters by size in $N \lg N$ time. The client can then read each cluster inventory item off the priority queue in order.

Aside from that, the bucket aspect of part 3 was fairly straightforward and used essentially the same architecture as part 2.

Part III Open Addressing

Part 3's open addressing section was the most difficult aspect of the entire project. It turns out that implementing an effective hashing method is very tricky, and I settled on fairly rudimentary functions for each of the supported keys. I experimented with functors - structs that expose the function call operators and rely upon function overloading for behavior specific to their types. This is how the c++ standard library seems to support plug-and-play coding, and I found it very effective.

I allow the client to specify three functors - a map capacity planner, a primary hash, and a secondary hash. The default map capacity planner takes the minimum capacity desired by the client, and returns the next-highest prime number to use as the true map capacity.

I was not able to get the `remove()` methods working for the quadratic probing or the double hashing instances. Quadratic probing tends to have a nasty habit of not reliably visiting every potential slot. Double hashing uses the key itself to pick the next slot, which leads to the following scenario:

- Keys A and B give equivalent primary hashing values but different double hashing values
- Key A exists in slot 1
- Key B is inserted, skipping slot 1 and probing forward to slot 2
- Key A is removed, and afterwards we attempt to resolve the cluster and reposition Key B so that it can be visited. However, Key A's probing value is dependent on Key A, so it is impossible to know that Key B originally intended to take Key A's slot, so when subsequently searching for Key B we will visit an empty slot where A was and encounter an empty slot

I'm sure there is a way to get the `remove()` methods working with double hashing and quadratic probing, but I ran out of time, and honestly I got really tired of fighting with the subtle mathematical landmine that hashing functions turns out to be. I have a new-found respect for cryptography researchers who deal with that kind of thing all the time.

Part IV

Part 4 was really fun. I implemented a BST base class and then exposed that to RBST and AVL. RBST simply overrode the insertion functionality to randomly insert at the root. Implementing all the methods for RBST was an exercise in recursive thinking, especially the pretty-print function which I was able to implement. I got a little carried away with playing with the pretty-print function, and I wanted to stress-test it to see how far I could push it. Because I allocate the output buffer (an array of lines) and write to it in memory before dumping it to standard output, a simple array requires contiguous memory and did not support the hundreds of millions of nodes that I wanted to print. Therefore, I used CDAL from the last project to store large chunks of lines together, but in different parts of memory.

As mentioned in the testing strategies, the BST class includes a few methods to manually and recursively calculate the height and number of children at each subtree if the `__DEBUG__` preprocessor macro is set to true. This is disabled by default because it is an expensive operation, but the unit tests set it to true, which makes ensuring the tree maintains correct structure throughout its lifetime.

Part IV Bonus: AVL Tree

The AVL tree was a bit tricky because I had to keep track of the height at every operation. I solved this once I went through and determined every operation that could potentially change the height of the subtree. Since each operation that could potentially change the height of a subtree was a recursive function, the base case would only change the height by a maximum of one. Therefore, at the base case, I call an `update_height` method, which simply calculates the height of the current node as one plus the maximum height between the two child nodes. After that operation completes, the recursion bubbles up, and the parent subtree does the same thing.

An elegant consequence of this strategy is that the AVL tree is simply the BST, but with the `insert()` and `remove()` functions wrapped in code that balances the current subtree at every recursion level. The `balance()` method is the largest block of code in the class, which ends up being relatively short and easy to read through and see what's going on. In addition, I wrap a few other BST methods with code that, if `__DEBUG__` is set to true, will recursively verify that `abs(balance factor)` is less than or equal to one at every subtree. Like the code that does the same for height and number of children, this is an expensive operation and is disabled by default, but it is very useful for unit testing.

Part I: Hashmap with Open Addressing

Part 1 Testing Strategy

Paul Nickerson

I attempted to separate tests into two somewhat distinct categories: operations that are expected to fail (`operation_failures.cpp`), and operations that are expected to succeed (`operation_successes.cpp`).

`operation_failures.cpp`

Within the failure operations scenario, I start with an empty map and try to `search()` and `remove()` an item whose key does not exist in the map. Both calls should return false. I then fill the map to capacity, which is possible to do since linear probing allows us to reliably fill every slot, and attempt to `insert()` a key (should fail due to lack of space), `remove()` a key that doesn't exist in the map, and `search()` for a key that does not exist in the map.

`operation_successes.cpp`

The success-expecting operations is much more extensive. I start by filling the map halfway, clearing it, then filling it up halfway again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the `print()` function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. I ensure that these match the value of `capacity() - size()`.

I then attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

part1/part1.pdf

part1/checklist.txt,

Hashmap with Open Addressing written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part I: hashmaps with Open Addressing
=====

My MAP implementation uses the data structure described in the part I
instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly supports the following methods
as described in part I:

- * insert: yes
- * remove: yes
- * search: yes
- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes

=====
FOR ALL PARTS
=====

My MAP implementation compiles correctly using g++ v4.8.2 on the
OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responses I have provided above are 100% true.
Should it be determined that any are not 100% true, I agree to take a 0
(zero) on the assignment: yes

I affirm that I am the sole author of this hashmaps with Open Addressing
and the associated tests.
Paul Nickerson, Dec 2 in COP3520 section 1087


```
How to compile and run my unit tests on the OpenBSD VM
cd part1/source
./compile.sh
./run_tests > output.txt
```

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29
30 namespace hash_utils {
31     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
32     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
33         choose a prime-number map capacity
34         251,
35         509,
36         1021,
37         2039,
38         4093,
39         8191,
40         16381,
41         32749,
42         65521,
43         131071,
44         262193,
45         524287,
46         1048573,
47         2097143,
```



```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```

```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
              number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
              loop).
148             //also, hash must be relatively prime to map_capacity so that
              every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
              so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178         for the types we need to support other than const char* (ie int,
179         double, and std::string),
180         we can pass these around willy-nilly. for const char*, handled
181         below, we will obtain our
182         own copy of the character array by wrapping it in a std::string
183     */
184 private:
185     T raw;
186     functors::compare compare;
187 public:
188     GenericContainer(const T& val): raw(val) {}
189     GenericContainer() = default;
190     GenericContainer& operator=(GenericContainer const& rhs) = delete;
191     T operator()() const {
192         return raw;
193     }

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```

```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```

```
342 }  
343  
344 #endif
```

common/priority_queue.h

common/priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12                 cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                 && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                     tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```



```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72 }
73 #endif // _PRIORITY_QUEUE_H_

```

part1/source/open_addressing_map.h

part1/source/open_addressing_map.h

```
1  #ifndef _OPEN_ADDRESSING_MAP_H_
2  #define _OPEN_ADDRESSING_MAP_H_
3
4  #include <iostream>
5  #include "../common/common.h"
6  #include <stdexcept>
7
8  namespace cop3530 {
9      class HashMapOpenAddressing {
10     private:
11         typedef int key_type;
12         typedef char value_type;
13         typedef hash_utils::ClusterInventory ClusterInventory;
14         struct Slot {
15             key_type key;
16             value_type value;
17             bool is_occupied = false;
18         };
19         Slot* slots;
20         size_t curr_capacity = 0;
21         size_t num_occupied_slots = 0;
22         size_t probe(size_t i) {
23             return i;
24         }
25         size_t hash(key_type const& key) {
26             size_t M = capacity();
27             hash_utils::functors::primary_hashes::hash_basic hasher;
28             size_t big_hash_number = hasher(key);
29             size_t hash_val = big_hash_number % M;
30             return hash_val;
31         }
32         /*
33          searches the map for an item matching key. returns the number of probe
34          attempts needed
35          to reach either the item or an empty slot
36          */
37         int search_internal(key_type const& key) {
38             size_t M = capacity();
39             size_t hash_val = hash(key);
40             size_t probe_index;
41             for (probe_index = 0; probe_index != M; ++probe_index) {
42                 size_t slot_index = (hash_val + probe(probe_index)) % M;
43                 if (slots[slot_index].is_occupied) {
44                     if (slots[slot_index].key == key) {
45                         //found the key
46                         break;
47                     }
48                 }
49             }
50         }
51     };
52 }
```

```

47         } else
48             //found unoccupied slot
49             break;
50     }
51     return 1 + probe_index; //start with a single probe when probe_index==0
52 }
53 //all backing array manipulations should go through the following two
54 //methods
55 void insert_at_index(key_type const& key, value_type const& value, size_t
56 index) {
57     Slot& s = slots[index];
58     s.key = key;
59     s.value = value;
60     if ( ! s.is_occupied) {
61         s.is_occupied = true;
62         ++num_occupied_slots;
63     }
64 }
65 value_type remove_at_index(size_t index) {
66     Slot& s = slots[index];
67     if (s.is_occupied) {
68         s.is_occupied = false;
69         --num_occupied_slots;
70     }
71     return s.value;
72 }
73 public:
74 HashMapOpenAddressing(size_t const min_capacity)
75 {
76     if (min_capacity == 0) {
77         throw std::domain_error("min_capacity must be at least 1");
78     }
79     cop3530::hash_utils::functors::map_capacity_planner capacity_planner;
80     curr_capacity = capacity_planner(min_capacity); //make capacity prime
81     slots = new Slot[curr_capacity];
82 }
83 ~HashMapOpenAddressing() {
84     delete slots;
85 }
86 /*
87     if there is space available, adds the specified key/value-pair to the
88     hash map and returns true; otherwise
89     returns false. If an item already exists in the map with the same key,
90     replace its value.
91 */
92 bool insert(key_type const& key, value_type const& value) {
93     size_t M = capacity();
94     if (M == size())
95         return false;
96     size_t probes_required = search_internal(key);
97     size_t index = (hash(key) + probe(probes_required - 1)) % M;
98     insert_at_index(key, value, index);

```

```

95         return true;
96     }
97     /*
98         if there is an item matching key, removes the key/value-pair from the
99         map, stores it's value in value,
100         and returns true; otherwise returns false.
101     */
102     bool remove(key_type const& key, value_type& value) {
103         size_t M = capacity();
104         size_t probes_required = search_internal(key);
105         size_t index = (hash(key) + probe(probes_required - 1)) % M;
106         if ( ! (slots[index].is_occupied && slots[index].key == key))
107             //key not found
108             return false;
109         value = remove_at_index(index);
110         size_t start_index = index;
111         //remove and reinsert items until find unoccupied slot
112         for (int i = 1; ; ++i) {
113             index = (start_index + probe(i)) % M;
114             Slot const& s = slots[index];
115             if (s.is_occupied) {
116                 remove_at_index(index);
117                 insert(s.key, s.value);
118             } else {
119                 break;
120             }
121         }
122         return true;
123     }
124     /*
125         if there is an item matching key, stores it's value in value,
126         and returns true (the item remains in the map); otherwise returns false.
127     */
128     bool search(key_type const& key, value_type& value) {
129         size_t M = capacity();
130         size_t probes_required = search_internal(key);
131         size_t index = (hash(key) + probe(probes_required - 1)) % M;
132         if ( ! (slots[index].is_occupied && slots[index].key == key))
133             //key not found
134             return false;
135         value = slots[index].value;
136         return true;
137     }
138     /*
139         removes all items from the map.
140     */
141     void clear() {
142         size_t cap = capacity();
143         for (size_t i = 0; i != cap; ++i)
144             slots[i].is_occupied = false;
145         num_occupied_slots = 0;
146     }

```

```

146     /*
147         returns true IFF the map contains no elements.
148     */
149     bool is_empty() {
150         return size() == 0;
151     }
152     /*
153         returns the number of slots in the map.
154     */
155     size_t capacity() {
156         return curr_capacity;
157     }
158     /*
159         returns the number of items actually stored in the map.
160     */
161     size_t size() {
162         return num_occupied_slots;
163     }
164     /*
165         returns the map's load factor (size = load * capacity).
166     */
167     double load() {
168         return static_cast<double>(size()) / capacity();
169     }
170     /*
171         inserts into the ostream, the backing array's contents in sequential
172         order.
173         Empty slots shall be denoted by a hyphen, non-empty slots by that item's
174         key. [This function will be used for debugging/monitoring].
175     */
176     std::ostream& print(std::ostream& out) {
177         size_t cap = capacity();
178         out << '[';
179         for (size_t i = 0; i != cap; ++i) {
180             if (slots[i].is_occupied) {
181                 out << slots[i].key;
182             } else {
183                 out << "-";
184             }
185             if (i + 1 < cap)
186                 out << '|';
187         }
188         out << ']';
189         return out;
190     }
191 };
192 }
193
194 #endif

```

Part II: Hashmap with Buckets

Part 2 Testing Strategy

Paul Nickerson

I attempted to separate tests into two somewhat distinct categories: operations that are expected to fail (`operation_failures.cpp`), and operations that are expected to succeed (`operation_successes.cpp`).

`operation_failures.cpp`

Within the failure operations scenario, I start with an empty map and try to `search()` and `remove()` an item whose key does not exist in the map. Both calls should return false. I then fill the map with a bunch of items (it is impossible to run out of space because collisions are resolved via an arbitrarily-growable linked list). From this newly-filled map, I attempt to `remove()` a key that doesn't exist in the map, and `search()` for a key that does not exist in the map, both of which should return false.

`operation_successes.cpp`

The success-expecting operations is much more extensive. I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the `print()` function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. Since $\text{load factor} = \text{occupied buckets} / \text{capacity}$, we can get the number of unoccupied buckets as $\text{capacity} * (1 - \text{load})$. This should equal the number of hyphens in the `print()` output.

I then attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

part2/part2.pdf

part2/checklist.txt

Hashmaps with Buckets written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part II: Hashmaps with Buckets
=====

My MAP implementation uses the data structure described in the part II instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly supports the following methods as described in part I:

- * insert: yes
- * remove: yes
- * search: yes
- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes

=====
FOR ALL PARTS
=====

My MAP implementation compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responses I have provided above are 100% true. Should it be determined that any are not 100% true, I agree to take a 0 (zero) on the assignment: yes

I affirm that I am the sole author of this Hashmaps with Buckets and the associated tests.
Paul Nickerson, Dec 2 in COP3520 section 1087


```
How to compile and run my unit tests on the OpenBSD VM
cd part2/source
./compile.sh
./run_tests > output.txt
```

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29
30 namespace hash_utils {
31     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
32     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
33         choose a prime-number map capacity
34         251,
35         509,
36         1021,
37         2039,
38         4093,
39         8191,
40         16381,
41         32749,
42         65521,
43         131071,
44         262193,
45         524287,
46         1048573,
47         2097143,
```

```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```

```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
              number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
              loop).
148             //also, hash must be relatively prime to map_capacity so that
              every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
              so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178         for the types we need to support other than const char* (ie int,
              double, and std::string),
179         we can pass these around willy-nilly. for const char*, handled
              below, we will obtain our
180         own copy of the character array by wrapping it in a std::string
181     */
182 private:
183     T raw;
184     functors::compare compare;
185 public:
186     GenericContainer(const T& val): raw(val) {}
187     GenericContainer() = default;
188     GenericContainer& operator=(GenericContainer const& rhs) = delete;
189     T operator()() const {
190         return raw;

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator>()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```

```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```



```
342 }  
343  
344 #endif
```

common/priority_queue.h

common/priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12             cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                 && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                     tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```

```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72
73 #endif // _PRIORITY_QUEUE_H_

```

part2/source/buckets_map.h

part2/source/buckets_map.h

```
1  #ifndef _BUCKETS_MAP_H_
2  #define _BUCKETS_MAP_H_
3
4  #include <iostream>
5  #include <stdexcept>
6  #include "../common/common.h"
7
8  namespace cop3530 {
9      class HashMapBuckets {
10     private:
11         typedef int key_type;
12         typedef char value_type;
13         typedef hash_utils::ClusterInventory ClusterInventory;
14         struct Item {
15             key_type key;
16             value_type value;
17             Item* next;
18             bool is_dummy;
19             Item(Item* next): next(next), is_dummy(true) {}
20         };
21         struct Bucket {
22             Item* head; //use a head pointer to the first node, and include a dummy
23                         //node at the end (but dont store its pointer)
24             Bucket() {
25                 Item* tail = new Item(nullptr);
26                 head = tail;
27             }
28             ~Bucket() {
29                 while ( ! head->is_dummy) {
30                     Item* to_delete = head;
31                     head = head->next;
32                     delete to_delete;
33                 }
34                 delete head; //tail
35             }
36         };
37         typedef Item* link;
38         Bucket* buckets;
39         size_t num_buckets = 0;
40         size_t num_items = 0;
41         size_t hash(key_type const& key) {
42             size_t M = capacity();
43             hash_utils::functors::primary_hashes::hash_basic hasher;
44             return hasher(key) % M;
45         }
46         /*
47          searches the bucket corresponding to the specified key's hash for that
```

```

47         key. if found, stores a reference to that item and returns P, the number
           of
48         probe attempts needed to get to the item (ie the number of chain links
           needed
49         to be traversed). otherwise return -1 * P and stores the pointer to the
           tail dummy node in
50         item_ptr.
51     */
52     int search_internal(key_type const& key, link& item_ptr) {
53         int probe_attempts = 1;
54         size_t hash_val = hash(key);
55         Bucket& bucket = buckets[hash_val];
56         item_ptr = bucket.head;
57         while ( ! item_ptr->is_dummy) {
58             if (item_ptr->key == key) {
59                 //found the key
60                 return probe_attempts;
61             }
62             item_ptr = item_ptr->next;
63             ++probe_attempts;
64         }
65         //key not found
66         return probe_attempts * -1;
67     }
68     void init() {
69         buckets = new Bucket[num_buckets];
70         num_items = 0;
71     }
72 public:
73     HashMapBuckets(size_t const min_buckets)
74     {
75         if (min_buckets == 0) {
76             throw std::domain_error("min_buckets must be at least 1");
77         }
78         cop3530::hash_utils::functors::map_capacity_planner capacity_planner;
79         num_buckets = capacity_planner(min_buckets); //make capacity prime
80         init();
81     }
82     ~HashMapBuckets() {
83         delete[] buckets;
84     }
85     /*
86         if there is space available, adds the specified key/value-pair to the
           hash map and returns true; otherwise
87         returns false. If an item already exists in the map with the same key,
           replace its value.
88         note: this will never return false because we add to a linked list to
           resolve collisions
89     */
90     bool insert(key_type const& key, value_type const& value) {
91         Item* item;
92         int probes_required = search_internal(key, item);

```

```

93         if (probes_required > 0)
94             //found item
95             item->value = value;
96         else {
97             //currently holding tail (item not found). transform it into a valid
98             item then add a new tail
99             item->is_dummy = false;
100             item->key = key;
101             item->value = value;
102             item->next = new Item(nullptr);
103             ++num_items;
104         }
105         return true;
106     }
107     /*
108     if there is an item matching key, removes the key/value-pair from the
109     map, stores it's
110     value in value, and returns true; otherwise returns false.
111     */
112     bool remove(key_type const& key, value_type& value) {
113         Item* item;
114         int probes_required = search_internal(key, item);
115         if (probes_required > 0) {
116             //found item
117             value = item->value;
118             //swap the current item for the next one
119             Item* to_delete = item->next;
120             *item = *to_delete;
121             delete to_delete;
122             --num_items;
123             return true;
124         }
125         return false;
126     }
127     /*
128     if there is an item matching key, stores it's value in value, and
129     returns true (the
130     item remains in the map); otherwise returns false.
131     */
132     bool search(key_type const& key, value_type& value) {
133         Item* item;
134         int probes_required = search_internal(key, item);
135         if (probes_required > 0) {
136             //found item
137             value = item->value;
138             return true;
139         }
140         return false;
141     }
142     /*
143     removes all items from the map.
144     */

```

```

142     void clear() {
143         delete[] buckets;
144         init();
145     }
146     /*
147         returns true IFF the map contains no elements.
148     */
149     bool is_empty() {
150         return size() == 0;
151     }
152     /*
153         returns the number of slots in the map.
154     */
155     size_t capacity() {
156         return num_buckets;
157     }
158     /*
159         returns the number of items actually stored in the map.
160     */
161     size_t size() {
162         return num_items;
163     }
164     /*
165         returns the map's load factor (occupied buckets = load * capacity).
166     */
167     double load() {
168         size_t occupied_buckets = 0;
169         if (size() > 0) {
170             size_t M = capacity();
171             for (size_t i = 0; i != M; ++i) {
172                 Bucket const& bucket = buckets[i];
173                 if ( ! bucket.head->is_dummy)
174                     //bucket has at least one item
175                     occupied_buckets++;
176             }
177         }
178         return static_cast<double>(occupied_buckets) / capacity();
179     }
180     /*
181         inserts into the ostream, the backing array's contents in sequential
182         order.
183         Empty slots shall be denoted by a hyphen, non-empty slots by that item's
184         key. [This function will be used for debugging/monitoring].
185     */
186     std::ostream& print(std::ostream& out) {
187         size_t cap = capacity();
188         bool print_separator = false;
189         out << '[';
190         for (size_t i = 0; i != cap; ++i) {
191             Bucket const& bucket = buckets[i];
192             if (bucket.head->is_dummy) {
193                 if (print_separator)

```

```

193         out << "|";
194     else
195         print_separator = true;
196         out << "-";
197     } else {
198         for (Item* item = bucket.head; item->is_dummy != true; item =
199             item->next) {
200             if (print_separator)
201                 out << "|";
202             else
203                 print_separator = true;
204             out << item->key;
205         }
206     }
207     out << ']' ;
208     return out;
209 }
210 };
211 }
212
213 #endif

```

Part III: Parameterizable Hashmap with Open Addressing

Part 3 Open Addressing Testing Strategy

Paul Nickerson

In addition to separating tests by operations that are expected to fail and those that are expected to succeed, I also included a scenario for testing the `cluster_distribution()` function.

operation__failures__(probe type/double hash).cpp

I include three failure operations scenarios, one for each of the secondary hashing methods to be supported (linear probing, quadratic probing, and double hashing). Within each scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string).

Linear Probing

To the linear probing instance, I start with an empty map and try to `search()` and `remove()` an item whose key does not exist in the map. Both calls should return a value less than zero (indicating key not found). I then fill the map to capacity, and attempt to insert an item where there is no space for it, which is expected to return a value less than zero. From this newly-filled map, I attempt to `remove()` a key that doesn't exist in the map, and `search()` for a key that does not exist in the map, both of which should return a value less than zero.

Quadratic Probing and Double Hashing

I was not able to get the `remove()` methods working for the quadratic probing or the double hashing instances. Quadratic probing tends to have a nasty habit of not reliably visiting every potential slot. Double hashing uses the key itself to pick the next slot, which leads to the following scenario:

- Keys A and B give equivalent primary hashing values but different double hashing values
- Key A exists in slot 1
- Key B is inserted, skipping slot 1 and probing forward to slot 2
- Key A is removed, and afterwards we attempt to resolve the cluster and reposition Key B so that it can be visited. However, Key A's probing value is dependent on Key A, so it is impossible to know that Key B originally intended to take Key A's slot, so

when subsequently searching for Key B we will visit an empty slot where A was and encounter an empty slot

Because of these issues, I did not test the `remove()` functionality of the quadratic probing and double hashing instances, opting instead to simply fill the map to capacity and then search for a key that should not exist in the map. Note that quadratic hashing, as previously mentioned, has a tendency to sometimes ignore some slots, and the probability of this occurring increases drastically as the load factor approaches 1. I early-abandon when this becomes evident to avoid an infinite loop, but, as a result, it was not possible to include a test to try and insert into a completely filled map like I could do in linear probing, since one key may induce an infinite loop while another key successfully finds one of the last remaining slots.

operation__successes__(probe type/double hash).cpp

Within the three success operations scenarios (one for each of the hashing methods), I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the `print()` function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. The number of hyphens in the `print()` output should equal `capacity() - size()`, ie the number of unoccupied slots.

As previously mentioned, I was unable to successfully implement `remove()` functionality in quadratic probing and double hashing, but within linear probing I attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

cluster__distribution__(probe type/double hash).cpp

Since `cluster_distribution()` returns a priority queue of clusters, and each cluster has a minimum size of one, all the clusters taken together should encompass every occupied slot. Therefore, I fill the map with a bunch of items, then clear it and fill it again to try and destabilize the map. From there, I take the summation, over every cluster, of the cluster's size times the number of clusters having that size. The result of that summation should equal the output from the map's `size()` method. I do this four times for each of the hashing method instances, once for each of the key types supported (a total of 12 times).

part3/open_addressing/part3open.pdf

part3/open_addressing/checklist.txt

hashmaps with Open Addressing written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part III: hashmaps with Open Addressing
=====

My MAP implementation uses the data structure described in the part II
instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly supports all three probing
techniques: (no - everything works except removing items with double hashing and quadratic probing)

My MAP implementation 100% correctly supports the following key types:

- * signed int: yes
- * double: yes
- * c-string: yes
- * std::string: yes

My MAP implementation 100% correctly supports the ANY value type: yes

My MAP implementation 100% correctly supports the following methods
as described in part I:

- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes

My MAP implementation 100% correctly supports the following revised
and new methods as described in part III:

- * insert: yes
- * remove: (linear probing: yes, quadratic probing: no, double hashing: no)
- * search: yes
- * cluster_distribution(): yes
- * remove_random(): (linear probing: yes, quadratic probing: no, double hashing: no)

=====
FOR ALL PARTS
=====

My MAP implementation compiles correctly using g++ v4.8.2 on the
OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responsess I have provided above are 100% true.
Should it be determined that any are not 100% true, I agree to take a 0
(zero) on the assignment: yes

I affirm that I am the sole author of this hashmaps with Open Addressing
and the associated tests.

Paul Nickerson, Dec 2 in COP3520 section 1087

How to compile and run my unit tests on the OpenBSD VM
cd part3/open_addressing/source
./compile.sh
./run_tests > output.txt

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29 }
30
31 namespace hash_utils {
32     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
33     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
34         choose a prime-number map capacity
35         251,
36         509,
37         1021,
38         2039,
39         4093,
40         8191,
41         16381,
42         32749,
43         65521,
44         131071,
45         262193,
46         524287,
47         1048573,
48         2097143,
```

```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```



```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
                number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
                loop).
148             //also, hash must be relatively prime to map_capacity so that
                every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
                so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178         for the types we need to support other than const char* (ie int,
            double, and std::string),
179         we can pass these around willy-nilly. for const char*, handled
            below, we will obtain our
180         own copy of the character array by wrapping it in a std::string
181     */
182 private:
183     T raw;
184     functors::compare compare;
185 public:
186     GenericContainer(const T& val): raw(val) {}
187     GenericContainer() = default;
188     GenericContainer& operator=(GenericContainer const& rhs) = delete;
189     T operator()() const {
190         return raw;

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             //deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```

```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```

```
342 }  
343  
344 #endif
```

common/priority_queue.h

common/priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12                 cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                 && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                     tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```

```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72
73 #endif // _PRIORITY_QUEUE_H_

```

part3/open_addressing/source/open_addressing_generic_map.h

part3/open_addressing/source/open_addressing_generic_map.h

```
1  #ifndef _HASHMAPOPENADDRESSINGGENERIC_H_
2  #define _HASHMAPOPENADDRESSINGGENERIC_H_
3
4  #include <iostream>
5  #include <string>
6  #include "../common/common.h"
7  #include "../common/priority_queue.h"
8
9  namespace cop3530 {
10     template<typename key_type,
11             typename value_type,
12             typename capacity_plan_functor =
13                 hash_utils::functors::map_capacity_planner,
14             typename primary_hash =
15                 hash_utils::functors::primary_hashes::hash_basic,
16             typename secondary_hash =
17                 hash_utils::functors::secondary_hashes::hash_double>
18     class HashMapOpenAddressingGeneric {
19     private:
20         typedef hash_utils::ClusterInventory ClusterInventory;
21         typedef hash_utils::Key<key_type, primary_hash, secondary_hash> Key;
22         typedef hash_utils::Value<value_type> Value;
23         struct Item {
24             Key key;
25             Value value;
26         };
27         struct Slot {
28             Item item;
29             bool is_occupied = false;
30         };
31         Slot* slots;
32         capacity_plan_functor choose_capacity;
33         size_t curr_capacity = 0;
34         size_t num_occupied_slots = 0;
35         /*
36          * searches the map for an item matching key. returns the number of probe
37          * attempts needed
38          * to reach either the item or an empty slot
39          */
40         int search_internal(Key const& key) {
41             size_t M = capacity();
42             size_t probe_index;
43             for (probe_index = 0; probe_index != M; ++probe_index) {
44                 size_t slot_index = key.hash(M, probe_index);
45                 if (slots[slot_index].is_occupied) {
46                     if (slots[slot_index].item.key == key) {
47                         //found the key
48                     }
49                 }
50             }
51         }
52     };
53 }
```



```

44         break;
45     }
46     } else
47         //found unoccupied slot
48         break;
49     }
50     return 1 + probe_index; //start with a single probe
51 }
52
53 //all backing array manipulations should go through the following two
54 //methods
55 void insert_at_index(Key const& key, Value const& value, size_t index) {
56     Slot& s = slots[index];
57     s.item.key.reset(key.raw_copy());
58     s.item.value.reset(value.raw_copy());
59     if ( ! s.is_occupied) {
60         s.is_occupied = true;
61         ++num_occupied_slots;
62     }
63 }
64 Value const& remove_at_index(size_t index) {
65     Slot& s = slots[index];
66     if (s.is_occupied) {
67         s.is_occupied = false;
68         --num_occupied_slots;
69     }
70     return s.item.value;
71 }
72 public:
73     HashMapOpenAddressingGeneric(size_t const min_capacity)
74     {
75         if (min_capacity == 0) {
76             throw std::domain_error("min_capacity must be at least 1");
77         }
78         curr_capacity = choose_capacity(min_capacity);
79         slots = new Slot[curr_capacity];
80     }
81     ~HashMapOpenAddressingGeneric() {
82         delete[] slots;
83     }
84
85     /*
86     if there is space available, adds the specified key/value-pair to the
87     hash map and returns the
88     number of probes required, P; otherwise returns -1 * P. If an item
89     already exists in the map
90     with the same key, replace its value.
91     */
92     int insert(key_type const& key, value_type const& value) {
93         size_t M = capacity();
94         if (M == size())
95             return -1 * size();

```

```

93     Key k(key);
94     Value v(value);
95     int probes_required = search_internal(k);
96     size_t index = k.hash(M, probes_required - 1);
97     if (slots[index].is_occupied == true && slots[index].item.key != k)
98         //map is full and we're going to hit an infinite loop if we keep
99         //going
100         return probes_required * -1;
101     insert_at_index(k, v, index);
102     return probes_required;
103 }
104
105 /*
106  * if there is an item matching key, removes the key/value-pair from the
107  * map, stores it's value in
108  * value, and returns the number of probes required, P; otherwise returns
109  * -1 * P.
110 */
111 int remove(key_type const& key, value_type& value) {
112     size_t M = capacity();
113     Key k(key);
114     int probes_required = search_internal(k);
115     size_t index = k.hash(M, probes_required - 1);
116     if (slots[index].is_occupied == false || slots[index].item.key != k)
117         //key not found
118         return -1 * probes_required;
119     Value v = remove_at_index(index);
120     value = v.raw_copy();
121     //remove and reinsert items until find unoccupied slot (guaranteed to
122     //happen since we just removed an item)
123     for (int i = 1; ; ++i) {
124         index = k.hash(M, i);
125         Slot const& s = slots[index];
126         if (s.is_occupied) {
127             remove_at_index(index);
128             insert(s.item.key.raw_copy(), s.item.value.raw_copy());
129         } else {
130             break;
131         }
132     }
133     return probes_required;
134 }
135
136 /*
137  * if there is an item matching key, stores it's value in value, and
138  * returns the
139  * number of probes required, P; otherwise returns -1 * P. Regardless, the
140  * item
141  * remains in the map.
142 */
143 int search(key_type const& key, value_type& value) {
144     size_t M = capacity();

```

```

139         Key k(key);
140         size_t probes_required = search_internal(k);
141         size_t index = k.hash(M, probes_required - 1);
142         if (slots[index].is_occupied == false || slots[index].item.key != k)
143             //key not found
144             return -1 * probes_required;
145         value = slots[index].item.value.raw_copy();
146         return probes_required;
147     }
148
149     /*
150     removes all items from the map.
151     */
152     void clear() {
153         size_t cap = capacity();
154         for (size_t i = 0; i != cap; ++i)
155             slots[i].is_occupied = false;
156         num_occupied_slots = 0;
157     }
158
159     /*
160     returns true IFF the map contains no elements.
161     */
162     bool is_empty() const {
163         return size() == 0;
164     }
165
166     /*
167     returns the number of slots in the map.
168     */
169     size_t capacity() const {
170         return curr_capacity;
171     }
172
173     /*
174     returns the number of items actually stored in the map.
175     */
176     size_t size() const {
177         return num_occupied_slots;
178     }
179
180     /*
181     returns the map's load factor (size = load * capacity).
182     */
183     double load() const {
184         return static_cast<double>(size()) / capacity();
185     }
186
187     /*
188     inserts into the ostream, the backing array's contents in sequential
189     order.
190     Empty slots shall be denoted by a hyphen, non-empty slots by that item's
191     key. [This function will be used for debugging/monitoring].
192     */
193     std::ostream& print(std::ostream& out) const {
194         size_t cap = capacity();
195         out << '[';

```

```

190     for (size_t i = 0; i != cap; ++i) {
191         if (slots[i].is_occupied) {
192             out << slots[i].item.key.raw();
193         } else {
194             out << "-";
195         }
196         if (i + 1 < cap)
197             out << '|';
198     }
199     out << ']' ;
200     return out;
201 }
202
203 priority_queue<ClusterInventory> cluster_distribution() {
204     //use an array to count cluster instances, then feed those to a priority
205     //queue and return it.
206     priority_queue<ClusterInventory> cluster_pq;
207     if (size() == 0) return cluster_pq;
208     size_t M = capacity();
209     size_t cluster_counter[M + 1];
210     for (size_t i = 0; i <= M; ++i)
211         cluster_counter[i] = 0;
212     if (size() == M) {
213         //handle the special case when the map is full
214         cluster_counter[size()]++;
215     } else {
216         //have at least one unoccupied slot
217         bool first_cluster_skipped = false;
218         size_t curr_cluster_size = 0;
219         //treat the backing array as a circular buffer and make a maximum of
220         //two passes to
221         //capture everything, including the wraparound cluster if it exists
222         for (size_t i = 1; i != M * 2; ++i) {
223             Slot const& curr_slot = slots[i % M], prev_slot = slots[(i - 1) %
224             M];
225             if (curr_slot.is_occupied && prev_slot.is_occupied) {
226                 //still in a cluster
227                 ++curr_cluster_size;
228             } else if (curr_slot.is_occupied && prev_slot.is_occupied ==
229             false) {
230                 //found a new cluster
231                 curr_cluster_size = 1;
232             } else if ( ! curr_slot.is_occupied && prev_slot.is_occupied) {
233                 //found the end of a cluster
234                 if (first_cluster_skipped) {
235                     cluster_counter[curr_cluster_size]++;
236                     if (i >= M) {
237                         //reached the end of the first cluster in the second
238                         //pass, so no all clusters have been handled
239                         break;
240                     }
241                 }
242             } else {

```

```

237         first_cluster_skipped = true;
238     }
239 }
240 }
241 }
242 for (size_t i = 1; i <= M; ++i)
243     if (cluster_counter[i] > 0) {
244         ClusterInventory cluster{i, cluster_counter[i]};
245         cluster_pq.add_to_queue(cluster);
246     }
247 return cluster_pq;
248 }
249
250 /*
251  generate a random number, R, (1,size), and starting with slot zero in
252  the backing array,
253  find the R-th occupied slot; remove the item from that slot (adjusting
254  subsequent items as
255  necessary), and return its key.
256 */
257 key_type remove_random() {
258     if (size() == 0) throw std::logic_error("Cant remove from an empty map");
259     size_t num_slots = capacity();
260     size_t ith_node_to_delete = 1 + rand_i(size());
261     for (size_t i = 0; i != num_slots; ++i) {
262         Slot const& slot = slots[i];
263         if (slot.is_occupied && --ith_node_to_delete == 0) {
264             key_type key = slot.item.key.raw_copy();
265             value_type val_dummy;
266             remove(key, val_dummy);
267             return key;
268         }
269     }
270     throw std::logic_error("Unexpected end of remove_random function");
271 }
272 };
273 #endif

```

Part III: Parameterizable Hashmap with Buckets

Part 3 Buckets Testing Strategy

Paul Nickerson

In addition to separating tests by operations that are expected to fail and those that are expected to succeed, I also included a scenario for testing the `cluster_distribution()` function.

operation_failures.cpp

Within the failure operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start with an empty map and try to `search()` and `remove()` an item whose key does not exist in the map. Both calls should return a value less than zero (indicating key not found). I then fill the map with a bunch of items (it is impossible to run out of space because collisions are resolved via an arbitrarily-growable linked list). From this newly-filled map, I attempt to `remove()` a key that doesn't exist in the map, and `search()` for a key that does not exist in the map, both of which should return a value less than zero.

operation_successes.cpp

Within the success operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the `print()` function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. Since $\text{load factor} = \text{occupied buckets} / \text{capacity}$, we can get the number of unoccupied buckets as $\text{capacity} * (1 - \text{load})$. This should equal the number of hyphens in the `print()` output.

I then attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

Part 3 added another function - `remove_random()` - which I test in a similar way to the preceding `remove()` check. I remove a random key, then try to `search()` for it and `remove()` it. Both checks should fail and return a value less than zero.

cluster_distribution.cpp

Since `cluster_distribution()` returns a priority queue of clusters, and each cluster has a minimum size of one, all the clusters taken together should encompass every occupied slot. Therefore, I fill the map with a bunch of items, then clear it and fill it again to try and destabilize the map. From there, I take the summation, over every cluster, of the cluster's size times the number of clusters having that size. The result of that summation should equal the output from the map's `size()` method. I do this four times, once for each of the four key types to be supported.

part3/bucket/part3bucket.pdf

part3/bucket/checklist.txt

Hashmaps with Buckets written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part III: Hashmaps with Buckets
=====

My MAP implementation uses the data structure described in the part II instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly supports the following key types:

- * signed int: yes
- * double: yes
- * c-string: yes
- * std::string: yes

My MAP implementation 100% correctly supports the ANY value type: yes

My MAP implementation 100% correctly supports the following methods as described in part I:

- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes

My MAP implementation 100% correctly supports the following revised and new methods as described in part III:

- * insert: yes
- * remove: yes
- * search: yes
- * cluster_distribution(): yes
- * remove_random(): yes

=====
FOR ALL PARTS
=====

My MAP implementation compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responsess I have provided above are 100% true.
Should it be determined that any are not 100% true, I agree to take a 0
(zero) on the assignment: yes

I affirm that I am the sole author of this Hashmaps with Buckets
and the associated tests.

Paul Nickerson, Dec 2 in COP3520 section 1087

How to compile and run my unit tests on the OpenBSD VM
cd part3/bucket/source
./compile.sh
./run_tests > output.txt

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29
30 namespace hash_utils {
31     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
32     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
33         choose a prime-number map capacity
34         251,
35         509,
36         1021,
37         2039,
38         4093,
39         8191,
40         16381,
41         32749,
42         65521,
43         131071,
44         262193,
45         524287,
46         1048573,
47         2097143,
```

```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```

```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
                number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
                loop).
148             //also, hash must be relatively prime to map_capacity so that
                every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
                so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178         for the types we need to support other than const char* (ie int,
            double, and std::string),
179         we can pass these around willy-nilly. for const char*, handled
            below, we will obtain our
180         own copy of the character array by wrapping it in a std::string
181     */
182 private:
183     T raw;
184     functors::compare compare;
185 public:
186     GenericContainer(const T& val): raw(val) {}
187     GenericContainer() = default;
188     GenericContainer& operator=(GenericContainer const& rhs) = delete;
189     T operator()() const {
190         return raw;

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             //deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```



```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```

```
342 }  
343  
344 #endif
```

common/priority_queue.h

common/priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12                 cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                    && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                                                  tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```

```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72
73 #endif // _PRIORITY_QUEUE_H_

```

part3/bucket/source/buckets_map.h

part3/bucket/source/buckets_map.h

```
1  #ifndef _BUCKETS_MAP_GENERIC_H_
2  #define _BUCKETS_MAP_GENERIC_H_
3
4  #include <iostream>
5  #include "../../common/common.h"
6  #include "../../common/SSL.h"
7  #include "../../common/priority_queue.h"
8
9  namespace cop3530 {
10     template<typename key_type,
11             typename value_type,
12             typename capacity_plan_funcor =
13                 hash_utils::functors::map_capacity_planner,
14             typename hash = hash_utils::functors::primary_hashes::hash_basic>
15     class HashMapBucketsGeneric {
16     private:
17         typedef hash_utils::ClusterInventory ClusterInventory;
18         typedef hash_utils::Key<key_type, hash> Key;
19         typedef hash_utils::Value<value_type> Value;
20         struct Item {
21             Key key;
22             Value value;
23             Item* next;
24             bool is_dummy;
25             explicit Item(Item* next): next(next), is_dummy(true) {}
26         };
27         struct Bucket {
28             Item* head; //use a head pointer to the first node, and include a dummy
29                         //node at the end (but dont store its pointer)
30             Bucket() {
31                 Item* tail = new Item(nullptr);
32                 head = tail;
33             }
34             ~Bucket() {
35                 while ( ! head->is_dummy) {
36                     Item* to_delete = head;
37                     head = head->next;
38                     delete to_delete;
39                 }
40                 delete head; //tail
41             }
42         };
43         typedef Item* link;
44         Bucket* buckets;
45         size_t num_buckets = 0;
46         size_t num_items = 0;
47         /*
```

```

46     searches the bucket corresponding to the specified key's hash for that
47     key. if found, stores a reference to that item and returns P, the number
        of
48     probe attempts needed to get to the item (ie the number of chain links
        needed
49     to be traversed). otherwise return -1 * P and stores the pointer to the
        tail dummy node in
50     item_ptr.
51 */
52 int search_internal(Key const& key, link& item_ptr) {
53     int probe_attempts = 1;
54     size_t hash_val = key.hash(capacity(), 0);
55     Bucket& bucket = buckets[hash_val];
56     item_ptr = bucket.head;
57     while ( ! item_ptr->is_dummy) {
58         if (item_ptr->key == key) {
59             //found the key
60             return probe_attempts;
61         }
62         item_ptr = item_ptr->next;
63         ++probe_attempts;
64     }
65     //key not found
66     return probe_attempts * -1;
67 }
68 void init() {
69     buckets = new Bucket[num_buckets];
70     num_items = 0;
71 }
72 public:
73     HashMapBucketsGeneric(size_t const min_buckets)
74     {
75         if (min_buckets == 0) {
76             throw std::domain_error("min_buckets must be at least 1");
77         }
78         cop3530::hash_utils::functors::map_capacity_planner capacity_planner;
79         num_buckets = capacity_planner(min_buckets); //make capacity prime
80         init();
81     }
82     ~HashMapBucketsGeneric() {
83         delete[] buckets;
84     }
85     /*
86     if there is space available, adds the specified key/value-pair to the
        hash map and returns the
87     number of probes required, P; otherwise returns -1 * P (that's a lie: we
        will always have space
88     available because each bucket contains a linked list that is
        indefinitely growable). If an item
89     already exists in the map with the same key, replace its value.
90 */
91 int insert(key_type const& key, value_type const& value) {

```

```

92     Item* item;
93     Key k(key);
94     Value v(value);
95     int probes_required = search_internal(k, item);
96     if (probes_required > 0)
97         //found item
98         item->value = v;
99     else {
100         //currently holding tail (item not found). transform it into a valid
            item then add a new tail
101         item->is_dummy = false;
102         item->key = k;
103         item->value = v;
104         item->next = new Item(nullptr);
105         ++num_items;
106     }
107     return std::abs(probes_required);
108 }
109 /*
110     if there is an item matching key, removes the key/value-pair from the
        map, stores it's value in
111     value, and returns the number of probes required, P; otherwise returns
        -1 * P.
112 */
113 int remove(key_type const& key, value_type& value) {
114     Key k(key);
115     Item* item;
116     int probes_required = search_internal(k, item);
117     if (probes_required > 0) {
118         //found item
119         value = item->value.raw_copy();
120         //swap the current item for the next one
121         Item* to_delete = item->next;
122         *item = *to_delete;
123         delete to_delete;
124         --num_items;
125     }
126     return probes_required;
127 }
128 /*
129     if there is an item matching key, stores it's value in value, and
        returns the
130     number of probes required, P; otherwise returns -1 * P. Regardless, the
        item
131     remains in the map.
132 */
133 int search(key_type const& key, value_type& value) {
134     Item* item;
135     Key k(key);
136     int probes_required = search_internal(k, item);
137     if (probes_required > 0) {
138         //found item

```

```

139         value = item->value.raw_copy();
140     }
141     return probes_required;
142 }
143 /*
144     removes all items from the map.
145 */
146 void clear() {
147     delete[] buckets;
148     init();
149 }
150 /*
151     returns true IFF the map contains no elements.
152 */
153 bool is_empty() {
154     return size() == 0;
155 }
156 /*
157     returns the number of slots in the map.
158 */
159 size_t capacity() {
160     return num_buckets;
161 }
162 /*
163     returns the number of items actually stored in the map.
164 */
165 size_t size() {
166     return num_items;
167 }
168 /*
169     returns the map's load factor (occupied buckets = load * capacity).
170 */
171 double load() {
172     size_t occupied_buckets = 0;
173     if (size() > 0) {
174         size_t M = capacity();
175         for (size_t i = 0; i != M; ++i) {
176             Bucket const& bucket = buckets[i];
177             if ( ! bucket.head->is_dummy)
178                 //bucket has at least one item
179                 occupied_buckets++;
180         }
181     }
182     return static_cast<double>(occupied_buckets) / capacity();
183 }
184 /*
185     inserts into the ostream, the backing array's contents in sequential
186     order.
187     Empty slots shall be denoted by a hyphen, non-empty slots by that item's
188     key. [This function will be used for debugging/monitoring].
189 */
190 std::ostream& print(std::ostream& out) {

```



```

190     size_t cap = capacity();
191     bool print_separator = false;
192     out << '[';
193     for (size_t i = 0; i != cap; ++i) {
194         Bucket const& bucket = buckets[i];
195         if (bucket.head->is_dummy) {
196             if (print_separator)
197                 out << "|";
198             else
199                 print_separator = true;
200             out << "-";
201         } else {
202             for (Item* item = bucket.head; item->is_dummy != true; item =
203                 item->next) {
204                 if (print_separator)
205                     out << "|";
206                 else
207                     print_separator = true;
208                 out << item->key.raw();
209             }
210         }
211     }
212     out << ']';
213     return out;
214 }
215
216 /*
217  returns a priority queue containing cluster sizes and instances (in the
218  form of ClusterInventory
219  struct instances), sorted by cluster size.
220 */
221 priority_queue<ClusterInventory> cluster_distribution() {
222     //use a simple linked list to count cluster instances, then feed those
223     //to a priority queue and return it.
224     priority_queue<ClusterInventory> cluster_pq;
225     if (size() == 0) return cluster_pq;
226     SSL<ClusterInventory> clusters;
227     size_t M = capacity();
228     for (size_t i = 0; i != M; ++i) {
229         Bucket const& bucket = buckets[i];
230         size_t bucket_size = 0;
231         Item* item_ptr = bucket.head;
232         while ( ! item_ptr->is_dummy) {
233             ++bucket_size;
234             item_ptr = item_ptr->next;
235         }
236         //I don't love this O(N^2) implementation, but premature
237         //optimization is the root of all evil and late projects
238         SSL<ClusterInventory>::iterator cluster_iterator = clusters.begin();
239         SSL<ClusterInventory>::iterator cluster_iterator_end =
240             clusters.end();
241         bool found_cluster = false;

```

```

237         for (; cluster_iterator != cluster_iterator_end; ++cluster_iterator)
238         {
239             if (cluster_iterator->cluster_size == bucket_size) {
240                 found_cluster = true;
241                 break;
242             }
243         }
244         if (found_cluster)
245             cluster_iterator->num_instances++;
246         else
247             clusters.push_back({bucket_size, 1});
248     }
249     SSL<ClusterInventory>::const_iterator cluster_iterator =
250         clusters.begin();
251     SSL<ClusterInventory>::const_iterator cluster_iterator_end =
252         clusters.end();
253     for (; cluster_iterator != cluster_iterator_end; ++cluster_iterator) {
254         if (cluster_iterator->cluster_size > 0)
255             cluster_pq.add_to_queue(*cluster_iterator);
256     }
257     return cluster_pq;
258 }
259
260 /*
261  generate a random number, R, (1,size), and starting with slot zero in
262  the backing array,
263  find the R-th occupied slot; remove the item from that slot (adjusting
264  subsequent items as
265  necessary), and return its key.
266 */
267 key_type remove_random() {
268     if (size() == 0) throw std::logic_error("Cant remove from an empty map");
269     size_t num_slots = capacity();
270     size_t ith_node_to_delete = 1 + rand_i(size());
271     for (size_t i = 0; i != num_slots; ++i) {
272         Bucket const& bucket = buckets[i];
273         Item* item_ptr = bucket.head;
274         while ( ! item_ptr->is_dummy) {
275             if (--ith_node_to_delete == 0) {
276                 key_type key = item_ptr->key.raw_copy();
277                 value_type val_dummy;
278                 remove(key, val_dummy);
279                 return key;
280             }
281             item_ptr = item_ptr->next;
282         }
283     }
284     throw std::logic_error("Unexpected end of remove_random function");
285 }
286 };
287
288 }
289

```

284 `#endif`

Part IV: Randomized BST

Part 4 RBST Testing Strategy

Paul Nickerson

operation_failures.cpp

Within the failure operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start with an empty map and try to search() and remove() an item whose key does not exist in the map. Both calls should return a value less than zero (indicating key not found). I then fill the map with a bunch of items (it is impossible to run out of space because collisions are resolved via an arbitrarily-growable linked list). From this newly-filled map, I attempt to remove() a key that doesn't exist in the map, and search() for a key that does not exist in the map, both of which should return a value less than zero.

The RBST class keeps track of the height and number of children of each subtree with $O(1)$ complexity during each operation that potentially changes those values. Because these values are so crucial to the map's functionality, and considering the difficulty of externally validating those values without adding more public methods, I include a preprocessor macro, `_DEBUG_`, which, when set to true (it defaults to false in production uses of the class), indicates to the BST base class to recursively verify the value of these values before each public method returns. This is potentially an expensive $O(N)$ operation that significantly slows down map functionality, which is why it is disabled by default and specifically enabled during unit testing. If it is determined that the $O(1)$ calculated values do not match the $O(N)$, "true" values, the class throws an exception, which is caught by the CATCH testing framework and causes the test to fail.

operation_successes.cpp

Within the success operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the print() function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. Since load factor = occupied buckets / capacity, we can get the number of unoccupied buckets as capacity * (1 - load). This should equal the number of hyphens in the print() output.

I then attempt to remove() several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both search() and remove() them, which should all return false.

I test `remove_random()` in a similar way to the preceding `remove()` check. I remove a random key, then try to `search()` for it and `remove()` it. Both checks should fail and return a value less than zero.

`tree__structure.cpp`

Because I was able to successfully implement the pretty-print bonus method, it is fairly straightforward to verify that the underlying RBST map successfully maintains the correct BST tree structure. I print the tree structure to a file, then make a `system()` call to pipe that through a few command line utilities that extract just the numerical key (I say numerical, but that number could be in the form of a string of `const char*`). This results in the natural order of the keys as they exist in the tree. I pipe those through `uniq` and `sort` to remove duplicates and induce expected sorting order. If the tree exhibits the correct structure, the output of those two operations will match verbatim.

I start by filling the map to half capacity, verify that `load()` returns 0.5, then check validate the tree structure. Afterwards, I fill the tree to full capacity and then delete half the nodes via `remove()` and `remove_random()` operations in an attempt to destabilize tree structure. Then I validate the tree structure again.

This series of checks is implemented for each of the four supported key types.

part4/part4.pdf

part4/checklist.txt

Randomized BST written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part IV: Randomized BST

=====
My MAP implementation uses the data structure described in the part IV
instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly implements RBST behavior: yes

My MAP implementation 100% correctly supports the following key types:

- * signed int: yes
- * double: yes
- * c-string: yes
- * std::string: yes

My MAP implementation 100% correctly supports the ANY value type: yes

My MAP implementation 100% correctly supports the following methods
as described in part IV:

- * insert: yes
- * remove: yes
- * search: yes
- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes
- * cluster_distribution(): yes
- * remove_random(): yes

My MAP implementation 100% correctly implements the bonus print(): yes

=====
FOR ALL PARTS

=====
My MAP implementation compiles correctly using g++ v4.8.2 on the
OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responses I have provided above are 100% true.
Should it be determined that any are not 100% true, I agree to take a 0
(zero) on the assignment: yes

I affirm that I am the sole author of this Randomized BST
and the associated tests.

Paul Nickerson, Dec 2 in COP3520 section 1087

How to compile and run my unit tests on the OpenBSD VM
cd part4/source
./compile.sh
./run_tests > output.txt

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29 }
30
31 namespace hash_utils {
32     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
33     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
34         choose a prime-number map capacity
35         251,
36         509,
37         1021,
38         2039,
39         4093,
40         8191,
41         16381,
42         32749,
43         65521,
44         131071,
45         262193,
46         524287,
47         1048573,
48         2097143,
```

```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```

```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
                number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
                loop).
148             //also, hash must be relatively prime to map_capacity so that
                every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
                so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178         for the types we need to support other than const char* (ie int,
            double, and std::string),
179         we can pass these around willy-nilly. for const char*, handled
            below, we will obtain our
180         own copy of the character array by wrapping it in a std::string
181     */
182 private:
183     T raw;
184     functors::compare compare;
185 public:
186     GenericContainer(const T& val): raw(val) {}
187     GenericContainer() = default;
188     GenericContainer& operator=(GenericContainer const& rhs) = delete;
189     T operator()() const {
190         return raw;

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator>()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             //deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```

```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```

```
342 }  
343  
344 #endif
```

common/priority_queue.h

common/priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12             cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                 && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                     tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```

```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72
73 #endif // _PRIORITY_QUEUE_H_

```

common/unit_test_utils.h

common/unit_test_utils.h

```
1  #ifndef _UNIT_TEST_UTILS_H_
2  #define _UNIT_TEST_UTILS_H_
3
4  #include <iostream>
5  #include <string>
6  #include <fstream>
7
8  namespace cop3530 {
9      namespace unit_test_utils {
10         inline std::string guid() {
11             std::string ret = "";
12             for (size_t i = 0; i != 32; ++i) {
13                 size_t rnd = cop3530::rand_i(16);
14                 if (rnd < 10)
15                     ret += std::string(1, '0' + rnd);
16                 else
17                     ret += std::string(1, 'A' + rnd - 10);
18             }
19             return ret;
20         }
21         inline std::string get_tmp_filename() {
22             return std::string("/tmp/") + guid() + std::string(".out");
23         }
24         inline void delete_file(std::string file_path) {
25             system((std::string("rm ") + file_path + std::string(" 2>&1 >>
26                 /tmp/debug")).c_str());
27         }
28         inline std::string shell_cmd(std::string cmd) {
29             std::string shell_script_file = get_tmp_filename();
30             std::string output_file = get_tmp_filename();
31             std::ofstream shell_script_out(shell_script_file);
32             shell_script_out << "#!/bin/sh" << std::endl << cmd << std::endl;
33             shell_script_out.close();
34             std::string chmod_cmd = std::string("chmod +x ") + shell_script_file;
35             system(chmod_cmd.c_str());
36             std::string invoke_cmd = shell_script_file + std::string(" > ") +
37                 output_file;
38             system(invoke_cmd.c_str());
39             std::ifstream read_output(output_file);
40             std::ostringstream oss;
41             std::string tmp;
42             while (std::getline(read_output, tmp)) {
43                 oss << tmp;
44                 if ( ! read_output.eof())
45                     oss << "\n";
46             }
47             delete_file(output_file);
48         }
49     }
50 }
```

```
46         delete_file(shell_script_file);
47         return oss.str();
48     }
49 }
50 }
51
52 #endif // _UNIT_TEST_UTILS_H_
```

part4/source/bst.h

part4/source/bst.h

```
1  #ifndef _BST_H_
2  #define _BST_H_
3
4  #include <cstdlib>
5  #include <sstream>
6  #include "../common/CDAL.h"
7  #include "../common/common.h"
8  #include "../common/priority_queue.h"
9
10 namespace cop3530 {
11     template<typename key_type,
12             typename value_type>
13     class BST {
14     protected: //let RBST and AVL inherit everything
15         typedef hash_utils::ClusterInventory ClusterInventory;
16         typedef hash_utils::Key<key_type> Key;
17         typedef hash_utils::Value<value_type> Value;
18         struct Node;
19         typedef Node* link;
20         struct Node {
21             Key key;
22             Value value;
23             size_t num_children;
24             size_t left_index;
25             size_t right_index;
26             size_t height; //height tracking coded in this class, but not used (for
                             //AVL, which is this class with self-balancing)
27             bool is_occupied;
28             size_t validate_children_count_recursive(Node* nodes) {
29                 //this function is for debugging purposes, does recursive traversal
                 //to find the correct number of children
30                 size_t child_count = 0;
31                 if (left_index)
32                     child_count += 1 +
                        nodes[left_index].validate_children_count_recursive(nodes);
33                 if (right_index)
34                     child_count += 1 +
                        nodes[right_index].validate_children_count_recursive(nodes);
35                 if (child_count != num_children) {
36                     std::ostringstream msg;
37                     msg << "Manually counted children, " << child_count << ",
                        different than child count, " << num_children;
38                     throw std::logic_error(msg.str());
39                 }
40                 return child_count;
41             }
42             size_t get_height_recursive(Node* nodes) {
```

```

43         //this function is for debugging purposes, does recursive traversal
           to find the correct height
44         size_t left_height = 0, right_height = 0;
45         size_t calculated_height = 0;
46         if (left_index)
47             left_height = nodes[left_index].get_height_recursive(nodes);
48         if (right_index)
49             right_height = nodes[right_index].get_height_recursive(nodes);
50         calculated_height = 1 + std::max(left_height, right_height);
51         return calculated_height;
52     }
53     void update_height(Node* nodes) {
54         //note: this method depends on the left and right subtree heights
           being correct
55         size_t left_height = 0, right_height = 0;
56         if (left_index)
57             left_height = nodes[left_index].height;
58         if (right_index)
59             right_height = nodes[right_index].height;
60         height = 1 + std::max(left_height, right_height);
61         if (_DEBUG_) {
62             size_t calculated_height = get_height_recursive(nodes);
63             if (calculated_height != height) {
64                 std::ostringstream msg;
65                 msg << "Manually calculated height, " << calculated_height <<
66                     ", different than tracked height, " << height;
67                 throw std::logic_error(msg.str());
68             }
69         }
70     void disable_and_adopt_free_tree(size_t free_index) {
71         is_occupied = false;
72         height = 0;
73         num_children = 0;
74         right_index = 0;
75         left_index = free_index;
76     }
77     void reset_and_enable(Key const& new_key, Value const& new_value) {
78         is_occupied = true;
79         height = 1; //self
80         left_index = right_index = 0;
81         num_children = 0;
82         key = new_key;
83         value = new_value;
84     }
85     int balance_factor(const Node* nodes) const {
86         size_t left_height = 0, right_height = 0;
87         if (left_index)
88             left_height = nodes[left_index].height;
89         if (right_index)
90             right_height = nodes[right_index].height;

```

```

91         return static_cast<long int>(left_height) - static_cast<long
           int>(right_height);
92     }
93 };
94 Node* nodes; /***note: array is 1-based so leaf nodes have child indices
           set to zero
95 size_t free_index;
96 size_t root_index;
97 size_t curr_capacity;
98 virtual size_t remove_smallest_key_node_index(size_t& subtree_root_index) {
99     //returns the index of the node with the smallest key, while
100     //setting its parent's left child index to the smallest key node's
101     //right child index. recursion downward through this function updates
102     //the heights of the nodes it traverses
103     Node& subtree_root = nodes[subtree_root_index];
104     size_t smallest_key_node_index = 0;
105     if (subtree_root_index == 0) {
106         throw std::logic_error("Expected to find a valid node, but didn't");
107     } else {
108         if (subtree_root.left_index) {
109             smallest_key_node_index =
110                 remove_smallest_key_node_index(subtree_root.left_index);
111             subtree_root.num_children--;
112             subtree_root.update_height(nodes);
113         } else {
114             smallest_key_node_index = subtree_root_index;
115             subtree_root_index = subtree_root.right_index;
116         }
117     }
118     return smallest_key_node_index;
119 }
120 virtual size_t remove_largest_key_node_index(size_t& subtree_root_index) {
121     //returns the index of the node with the largest key, while
122     //setting its parent's right child index to the largest key node's
123     //left child index. recursion downward through this function updates
124     //the heights of the nodes it traverses
125     Node& subtree_root = nodes[subtree_root_index];
126     size_t largest_key_node_index = 0;
127     if (subtree_root_index == 0) {
128         throw std::logic_error("Expected to find a valid node, but didn't");
129     } else {
130         if (subtree_root.right_index) {
131             largest_key_node_index =
132                 remove_largest_key_node_index(subtree_root.right_index);
133             subtree_root.num_children--;
134             subtree_root.update_height(nodes);
135         } else {
136             largest_key_node_index = subtree_root_index;
137             subtree_root_index = subtree_root.left_index;
138         }
139     }
140     return largest_key_node_index;

```

```

139     }
140     virtual void remove_node(size_t& subtree_root_index) {
141         Node& subtree_root = nodes[subtree_root_index];
142         size_t index_to_delete = subtree_root_index;
143         if (subtree_root.right_index || subtree_root.left_index) {
144             //subtree has at least one child
145             if (subtree_root.right_index)
146                 //replace the root with the smallest-keyed node in the right
147                 subtree
148                 subtree_root_index =
149                     remove_smallest_key_node_index(subtree_root.right_index);
150             else if (subtree_root.left_index)
151                 //replace the root with the largest-keyed node in the left subtree
152                 subtree_root_index =
153                     remove_largest_key_node_index(subtree_root.left_index);
154             //have the new root adopt the old root's children
155             Node& new_root = nodes[subtree_root_index];
156             new_root.left_index = subtree_root.left_index;
157             new_root.right_index = subtree_root.right_index;
158             //the new root has the same number of children as the old root,
159             //minus one
160             new_root.num_children = subtree_root.num_children - 1;
161             //removing the smallest/largest-keyed node from the old root has the
162             //effect of
163             //updating the heights of the old root's relevant subtrees (which
164             //the new root
165             //just adopted), so we can update the new root's height now
166             new_root.update_height(nodes);
167         } else
168             //neither subtree exists, so just delete the node
169             subtree_root_index = 0;
170         //node has been disowned by all ancestors, and has disowned all
171         //descendents, so free it
172         add_node_to_free_tree(index_to_delete);
173     }
174     virtual int do_remove(size_t nodes_visited, //starts at 0 when this
175                           function is first called (ie does not include current node visitation)
176                           size_t& subtree_root_index,
177                           Key const& key,
178                           Value& value,
179                           bool& found_key)
180     {
181         if (subtree_root_index != 0) {
182             Node& subtree_root = nodes[subtree_root_index];
183             ++nodes_visited;
184             //keep going down to the base of the tree
185             if (key < subtree_root.key) {
186                 nodes_visited = do_remove(nodes_visited, subtree_root.left_index,
187                                           key, value, found_key);
188             }
189             if (found_key) {
190                 //found the desired node and delete it
191                 subtree_root.num_children--;
192             }
193         }
194     }

```



```

182         //left child changed, so recompute subtree height
183         subtree_root.update_height(nodes);
184     }
185     } else if (key > subtree_root.key) {
186         nodes_visited = do_remove(nodes_visited,
187             subtree_root.right_index, key, value, found_key);
188         if (found_key) {
189             //found the desired node and delete it
190             subtree_root.num_children--;
191             //right child changed, so recompute subtree height
192             subtree_root.update_height(nodes);
193         }
194     } else if (key == subtree_root.key) {
195         //found key, remove the node
196         found_key = true;
197         value = subtree_root.value;
198         remove_node(subtree_root_index);
199     } else {
200         throw std::logic_error("Unexpected compare result");
201     }
202     }
203     return nodes_visited;
204 }
205 void write_subtree_buffer(size_t subtree_root_index,
206     CDAL<std::string>& buffer_lines,
207     size_t root_line_index,
208     size_t lbound_line_index /*inclusive*/,
209     size_t ubound_line_index /*exclusive*/) const
210 {
211     Node subtree_root = nodes[subtree_root_index];
212     std::ostringstream oss;
213     //print the node
214     //todo: fix this to only print the key
215     oss << "[" << subtree_root.key.raw() << "]";
216     buffer_lines[root_line_index] += oss.str();
217     //print the right descendents
218     if (subtree_root.right_index > 0) {
219         //at least 1 right child
220         size_t top_dashes = 1;
221         Node const& right_child = nodes[subtree_root.right_index];
222         if (right_child.left_index > 0) {
223             //right child has at least 1 left child
224             Node const& right_left_child = nodes[right_child.left_index];
225             top_dashes += 2 * (1 + right_left_child.num_children);
226         }
227         size_t top_line_index = root_line_index - 1;
228         while (top_line_index >= root_line_index - top_dashes)
229             buffer_lines[top_line_index--] += "| ";
230         size_t right_child_line_index = top_line_index;
231         buffer_lines[top_line_index--] += "+--";
232         while (top_line_index >= lbound_line_index)
233             buffer_lines[top_line_index--] += " ";

```

```

233         write_subtree_buffer(subtree_root.right_index,
234                               buffer_lines,
235                               right_child_line_index,
236                               lbound_line_index,
237                               root_line_index);
238     }
239     //print the left descendents
240     if (subtree_root.left_index > 0) {
241         //at least 1 left child
242         size_t bottom_dashes = 1;
243         Node const& left_child = nodes[subtree_root.left_index];
244         if (left_child.right_index > 0) {
245             //left child has at least 1 right child
246             Node const& left_right_child = nodes[left_child.right_index];
247             bottom_dashes += 2 * (1 + left_right_child.num_children);
248         }
249         size_t bottom_line_index = root_line_index + 1;
250         while (bottom_line_index <= root_line_index + bottom_dashes)
251             buffer_lines[bottom_line_index++] += "| ";
252         size_t left_child_line_index = bottom_line_index;
253         buffer_lines[bottom_line_index++] += "+--";
254         while (bottom_line_index < ubound_line_index)
255             buffer_lines[bottom_line_index++] += " ";
256         write_subtree_buffer(subtree_root.left_index,
257                               buffer_lines,
258                               left_child_line_index,
259                               root_line_index + 1,
260                               ubound_line_index);
261     }
262 }
263 void add_node_to_free_tree(size_t node_index) {
264     nodes[node_index].disable_and_adopt_free_tree(free_index);
265     free_index = node_index;
266 }
267 size_t procure_node(Key const& key, Value const& value) {
268     //updates the free index to the first free node's left child (while
269     //transforming that first free
270     //node to an enabled node with the specified key/value) and returns the
271     //index of what was the last
272     //free index
273     size_t node_index = free_index;
274     free_index = nodes[free_index].left_index;
275     Node& n = nodes[node_index];
276     n.reset_and_enable(key, value);
277     return node_index;
278 }
279 virtual int insert_at_leaf(size_t nodes_visited, //starts at 0 when this
280                           function is first called (ie does not include current node visitation)
281                           size_t& subtree_root_index,
282                           Key const& key,
283                           Value const& value,
284                           bool& found_key)

```

```

282     {
283         if (subtree_root_index == 0) {
284             //key not found
285             subtree_root_index = procure_node(key, value);
286         } else {
287             //parent was not a leaf
288             //keep going down to the base of the tree
289             Node& subtree_root = nodes[subtree_root_index];
290             ++nodes_visited;
291             if (key < subtree_root.key) {
292                 nodes_visited = insert_at_leaf(nodes_visited,
293                     subtree_root.left_index, key, value, found_key);
294                 if ( ! found_key) {
295                     //given key is unique to the tree, so a new node was added
296                     subtree_root.num_children++;
297                     subtree_root.update_height(nodes);
298                 }
299             } else if (key > subtree_root.key) {
300                 nodes_visited = insert_at_leaf(nodes_visited,
301                     subtree_root.right_index, key, value, found_key);
302                 if ( ! found_key) {
303                     //given key is unique to the tree, so a new node was added
304                     subtree_root.num_children++;
305                     subtree_root.update_height(nodes);
306                 }
307             } else if (key == subtree_root.key) {
308                 //found key, replace the value
309                 subtree_root.value = value;
310                 found_key = true;
311             } else {
312                 throw std::logic_error("Unexpected compare result");
313             }
314         }
315         return nodes_visited;
316     }
317 }
318
319 void rotate_left(size_t& subtree_root_index) {
320     Node& subtree_root = nodes[subtree_root_index];
321     size_t right_child_index = subtree_root.right_index;
322     Node& right_child = nodes[right_child_index];
323
324     //original root adopts the right child's left subtree
325     subtree_root.right_index = right_child.left_index;
326     //original root adopted a subtree (whose height did not change), so
327     //update its height
328     subtree_root.update_height(nodes);
329
330     //right child adopts original root and its children
331     right_child.left_index = subtree_root_index;
332     //right child (new root) adopted the original root (whose height has
333     //been updated), so update its height
334     right_child.update_height(nodes);

```

```

329         //since right child took the subtree root's place, it has the same
           number of children as the original root
330         right_child.num_children = subtree_root.num_children;
331
332         //root has new children, so update that counter (done after changing the
           right child's children counter
333         //because that depends on the original root's counter)
334         subtree_root.num_children = 0;
335         if (subtree_root.left_index != 0)
336             subtree_root.num_children += 1 +
           nodes[subtree_root.left_index].num_children;
337         if (subtree_root.right_index != 0)
338             subtree_root.num_children += 1 +
           nodes[subtree_root.right_index].num_children;
339
340         //set the right child as the new root
341         subtree_root_index = right_child_index;
342     }
343     void rotate_right(size_t& subtree_root_index) {
344         Node& subtree_root = nodes[subtree_root_index];
345         size_t left_child_index = subtree_root.left_index;
346         Node& left_child = nodes[left_child_index];
347
348         //original root adopts the left child's right subtree
349         subtree_root.left_index = left_child.right_index;
350         //original root adopted a subtree (whose height did not change), so
           update its height
351         subtree_root.update_height(nodes);
352
353         //left child adopts original root and its children
354         left_child.right_index = subtree_root_index;
355         //left child (new root) adopted the original root (whose height has been
           updated), so update its height
356         left_child.update_height(nodes);
357         //since left child took the subtree root's place, it has the same number
           of children as the original root
358         left_child.num_children = subtree_root.num_children;
359
360         //root has new children, so update that counter (done after changing the
           left child's children counter
361         //because that depends on the original root's counter)
362         subtree_root.num_children = 0;
363         if (subtree_root.left_index != 0)
364             subtree_root.num_children += 1 +
           nodes[subtree_root.left_index].num_children;
365         if (subtree_root.right_index != 0)
366             subtree_root.num_children += 1 +
           nodes[subtree_root.right_index].num_children;
367
368         //set the left child as the new root
369         subtree_root_index = left_child_index;
370     }

```

```

371 int do_search(size_t nodes_visited, //starts at 0 when this function is
    first called (ie does not include current node visitation)
372             size_t subtree_root_index,
373             Key const& key,
374             Value& value,
375             bool& found_key) const
376 {
377     if (subtree_root_index != 0) {
378         Node const& subtree_root = nodes[subtree_root_index];
379         ++nodes_visited;
380         if (key < subtree_root.key) {
381             nodes_visited = do_search(nodes_visited, subtree_root.left_index,
    key, value, found_key);
382         } else if (key > subtree_root.key) {
383             nodes_visited = do_search(nodes_visited,
    subtree_root.right_index, key, value, found_key);
384         } else if (key == subtree_root.key) {
385             //found key, replace the value
386             value = subtree_root.value;
387             found_key = true;
388         } else {
389             throw std::logic_error("Unexpected compare result");
390         }
391     }
392     return nodes_visited;
393 }
394 void prepare_cluster_distribution(size_t subtree_root_index,
    size_t curr_height, //includes the height of
    the current node, ie assumes current node
    exists
    size_t cluster_counter[])
396 {
397     Node const& subtree_root = nodes[subtree_root_index];
398     if ( ! subtree_root.left_index && ! subtree_root.right_index)
399         //at a leaf node
400         cluster_counter[curr_height]++;
401     else {
402         if (subtree_root.left_index)
403             prepare_cluster_distribution(subtree_root.left_index, curr_height
    + 1, cluster_counter);
404         if (subtree_root.right_index)
405             prepare_cluster_distribution(subtree_root.right_index,
    curr_height + 1, cluster_counter);
406     }
407 }
408
409 void remove_ith_node_inorder(size_t& subtree_root_index,
    size_t& ith_node_to_delete,
    Key& key)
410 {
411     Node& subtree_root = nodes[subtree_root_index];
412     if (subtree_root.left_index)
413

```

```

416         remove_ith_node_inorder(subtree_root.left_index, ith_node_to_delete,
417                                 key);
418     if (ith_node_to_delete == 0)
419         //deleted node in child subtree; nothing more to do
420         return;
421     if (--ith_node_to_delete == 0) {
422         //delete the current node
423         value_type dummy_val;
424         remove(subtree_root.key.raw_copy(), dummy_val);
425         key = subtree_root.key;
426         return;
427     }
428     if (subtree_root.right_index)
429         remove_ith_node_inorder(subtree_root.right_index,
430                                 ith_node_to_delete, key);
431 }
432
433 public:
434     /*
435     The constructor will allocate an array of capacity (binary
436     tree) nodes. Then make a chain from all the nodes (e.g.,
437     make node 2 the left child of node 1, make node 3 the left
438     child of node 2, &c. this is the initial free list.
439     */
440     BST(size_t capacity):
441         curr_capacity(capacity)
442     {
443         if (capacity == 0) {
444             throw std::domain_error("capacity must be at least 1");
445         }
446         nodes = new Node[capacity + 1];
447         clear();
448     }
449     /*
450     if there is space available, adds the specified key/value-pair to the
451     tree
452     and returns the number of nodes visited, V; otherwise returns -1 * V. If
453     an
454     item already exists in the tree with the same key, replace its value.
455     */
456     virtual int insert(key_type const& key, value_type const& value) {
457         if (size() == capacity())
458             //no more space
459             return -1 * size();
460         bool found_key = false;
461         Key k(key);
462         Value v(value);
463         int nodes_visited = insert_at_leaf(0, root_index, k, v, found_key);
464         if (_DEBUG_)
465             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
466         return nodes_visited;
467     }

```

```

464     /*
465         if there is an item matching key, removes the key/value-pair from the
            tree, stores
466         it's value in value, and returns the number of probes required, V;
            otherwise returns -1 * V.
467     */
468     virtual int remove(key_type const& key, value_type& value) {
469         if (is_empty())
470             return 0;
471         bool found_key = false;
472         Key k(key);
473         Value v(value);
474         int nodes_visited = do_remove(0, root_index, k, v, found_key);
475         if (_DEBUG_)
476             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
477         if (found_key)
478             value = v.raw_copy();
479         return found_key ? nodes_visited : -1 * nodes_visited;
480     }
481     /*
482         if there is an item matching key, stores it's value in value, and
            returns the number
483         of nodes visited, V; otherwise returns -1 * V. Regardless, the item
            remains in the tree.
484     */
485     virtual int search(key_type const& key, value_type& value) {
486         if (is_empty())
487             return 0;
488         bool found_key = false;
489         Key k(key);
490         Value v(value);
491         int nodes_visited = do_search(0, root_index, k, v, found_key);
492         if (found_key)
493             value = v.raw_copy();
494         return found_key ? nodes_visited : -1 * nodes_visited;
495     }
496     /*
497         removes all items from the map
498     */
499     virtual void clear() {
500         //Since I use size_t to hold the node indices, I make the node array
501         //1-based, with child index of 0 indicating that the current node is a
            leaf
502         for (size_t i = 1; i != capacity(); ++i)
503             nodes[i].disable_and_adopt_free_tree(i + 1);
504         free_index = 1;
505         root_index = 0;
506     }
507     /*
508         returns true IFF the map contains no elements.
509     */
510     virtual bool is_empty() const {

```

```

511         return size() == 0;
512     }
513     /*
514     returns the number of slots in the backing array.
515     */
516     virtual size_t capacity() const {
517         return curr_capacity;
518     }
519     /*
520     returns the number of items actually stored in the tree.
521     */
522     virtual size_t size() const {
523         if (root_index == 0) return 0;
524         Node const& root = nodes[root_index];
525         return 1 + root.num_children;
526     }
527     /*
528     [not a regular BST operation, but specific to this implementation]
529     returns the tree's load factor: load = size / capacity.
530     */
531     virtual double load() const {
532         return static_cast<double>(size()) / capacity();
533     }
534     /*
535     prints the tree in the following format:
536     +--[tiger]
537     | |
538     | | +--[panther]
539     | | |
540     | +--[ocelot]
541     | |
542     | +--[lion]
543     |
544     [leopard]
545     |
546     | +--[house cat]
547     | |
548     | +--[cougar]
549     | |
550     +--[cheetah]
551     |
552     +--[bobcat]
553     */
554     virtual std::ostream& print(std::ostream& out) const {
555         if (is_empty())
556             return out;
557         size_t num_lines = size() * 2 - 1;
558         //use CDAL here so we can print really super-huge trees where the write
559         //buffer doesn't fit in memory
560         CDAL<std::string> buffer_lines(100000);
561         for(size_t i = 0; i <= num_lines; ++i)
562             buffer_lines.push_back("");

```



```

562     Node const& root = nodes[root_index];
563     size_t root_line_index = 1;
564     if (root.right_index) {
565         root_line_index += 2 * (1 + nodes[root.right_index].num_children);
566     }
567     write_subtree_buffer(root_index, buffer_lines, root_line_index, 1,
568                          num_lines + 1);
569     for (size_t i = 1; i <= num_lines; ++i)
570         out << buffer_lines[i] << std::endl;
571     return out;
572 }
573
574 /*
575  returns a list indicating the number of leaf nodes at each height (since
576  the RBST doesn't exhibit
577  true clustering, but can have degenerate branches).
578 */
579 virtual priority_queue<hash_utils::ClusterInventory> cluster_distribution()
580 {
581     //use an array to count cluster instances, then feed those to a priority
582     queue and return it.
583     priority_queue<ClusterInventory> cluster_pq;
584     if (is_empty()) return cluster_pq;
585     size_t max_height = nodes[root_index].height;
586     size_t cluster_counter[max_height + 1];
587     for (size_t i = 0; i <= max_height; ++i)
588         cluster_counter[i] = 0;
589     prepare_cluster_distribution(root_index, 1, cluster_counter);
590     for (size_t i = 1; i <= max_height; ++i)
591         if (cluster_counter[i] > 0) {
592             ClusterInventory cluster{i, cluster_counter[i]};
593             cluster_pq.add_to_queue(cluster);
594         }
595     return cluster_pq;
596 }
597
598 /*
599  generate a random number, R, (1,size), and starting with the root (node
600  1), do an in-order
601  traversal to find the R-th occupied node; remove that node (adjusting
602  its children accordingly),
603  and return its key.
604
605  ***XXX: this likely contains a bug when using const char* keys in that
606  we'll be returning a dangling pointer!!!**
607 */
608 virtual key_type remove_random() {
609     if (size() == 0) throw std::logic_error("Cant remove from an empty map");
610     size_t ith_node_to_delete = 1 + rand_i(size());
611     Key key;
612     remove_ith_node_inorder(root_index, ith_node_to_delete, key);
613     key_type ret = key.raw_copy();

```

```
607         return ret;
608     }
609 };
610 }
611
612 #endif
```

part4/source/rbst.h

part4/source/rbst.h

```
1  #ifndef _RBST_H_
2  #define _RBST_H_
3
4
5  #include <cstdlib>
6  #include <sstream>
7  #include "../common/CDAL.h"
8  #include "../common/common.h"
9  #include "../common/priority_queue.h"
10 #include "bst.h"
11
12 namespace cop3530 {
13     template<typename key_type,
14             typename value_type>
15     class RBST: public BST<key_type, value_type> {
16     /*
17         Within the RBST insert_at_leaf method, the recursive execution path is
18         randomly redirected
19         to insert at the root. Therefore, we simply inherit from a generic BST
20         class and wrap the
21         insert_at_leaf method with that potential alternative execution path
22     */
23     private:
24         using super = BST<key_type, value_type>;
25         using typename super::Node;
26         typedef hash_utils::Key<key_type> Key;
27         typedef hash_utils::Value<value_type> Value;
28         int insert_at_leaf(size_t nodes_visited, //starts at 0 when this function
29                             is first called (ie does not include current node visitation)
30                             size_t& subtree_root_index,
31                             Key const& key,
32                             Value const& value,
33                             bool& found_key)
34         {
35             //parent was not a leaf
36             Node& subtree_root = this->nodes[subtree_root_index];
37             if (rand() < RAND_MAX / (subtree_root.num_children + 1)) {
38                 //randomly insert at the subtree root
39                 nodes_visited = insert_at_root(nodes_visited, subtree_root_index,
40                                                 key, value, found_key);
41             } else {
42                 nodes_visited = super::insert_at_leaf(nodes_visited,
43                                                         subtree_root_index, key, value, found_key);
44             }
45             return nodes_visited;
46         }
47     }
48     int insert_at_root(size_t nodes_visited,
```

```

43         size_t& subtree_root_index,
44         Key const& key,
45         Value const& value,
46         bool& found_key)
47     {
48         if (subtree_root_index == 0) {
49             //parent was a leaf, so create a new leaf
50             subtree_root_index = this->procure_node(key, value);
51         } else {
52             //parent was not a leaf
53             Node& subtree_root = this->nodes[subtree_root_index];
54             ++nodes_visited;
55             //keep going down to the base of the tree
56
57             if (key < subtree_root.key) {
58                 nodes_visited = insert_at_root(nodes_visited,
59                     subtree_root.left_index, key, value, found_key);
60                 if ( ! found_key) {
61                     //new node currently a new child of subtree root, so increment
62                     //the subtree root's number of children before rotating - new
63                     node
64                     //will adopt the root and its children and will take on the
65                     value
66                     //of its num_children
67                     subtree_root.num_children++;
68
69                     //current subtree root may have had its height changed, so
70                     update that before
71                     //promoting the new node
72                     subtree_root.update_height(this->nodes);
73                     this->rotate_right(subtree_root_index);
74                 }
75             } else if (key > subtree_root.key) {
76                 nodes_visited = insert_at_root(nodes_visited,
77                     subtree_root.right_index, key, value, found_key);
78                 if ( ! found_key) {
79                     subtree_root.num_children++;
80
81                     //current subtree root may have had its height changed, so
82                     update that before
83                     //promoting the new node
84                     subtree_root.update_height(this->nodes);
85                     this->rotate_left(subtree_root_index);
86                 }
87             } else if (key == subtree_root.key) {
88                 //found key, replace the value
89                 subtree_root.value = value;
90                 found_key = true;
91             } else {
92                 throw std::logic_error("Unexpected compare result");
93             }
94         }
95     }

```

```

89         return nodes_visited;
90     }
91 public:
92     RBST(size_t capacity): super(capacity) {}
93     /*
94         if there is space available, adds the specified key/value-pair to the
95         tree
96         and returns the number of nodes visited, V; otherwise returns -1 * V. If
97         an
98         item already exists in the tree with the same key, replace its value.
99     */
100     int insert(key_type const& key, value_type const& value) {
101         if (this->size() == this->capacity())
102             //no more space
103             return -1 * this->size();
104         bool found_key = false;
105         Key k(key);
106         Value v(value);
107         int nodes_visited = insert_at_leaf(0, this->root_index, k, v, found_key);
108         if (_DEBUG_)
109             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
110         return nodes_visited;
111     }
112 };
113 #endif

```

Part IV BONUS: AVL Tree

Part 4 AVL Testing Strategy

Paul Nickerson

operation_failures.cpp

Within the failure operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start with an empty map and try to search() and remove() an item whose key does not exist in the map. Both calls should return a value less than zero (indicating key not found). I then fill the map with a bunch of items (it is impossible to run out of space because collisions are resolved via an arbitrarily-growable linked list). From this newly-filled map, I attempt to remove() a key that doesn't exist in the map, and search() for a key that does not exist in the map, both of which should return a value less than zero.

The AVL class keeps track of the height and number of children of each subtree with $O(1)$ complexity during each operation that potentially changes those values. Because these values are so crucial to the map's functionality, and considering the difficulty of externally validating those values without adding more public methods, I include a preprocessor macro, `__DEBUG__`, which, when set to true (it defaults to false in production uses of the class), indicates to the BST base class to recursively verify the value of these values before each public method returns. This is potentially an expensive $O(N)$ operation that significantly slows down map functionality, which is why it is disabled by default and specifically enabled during unit testing. If it is determined that the $O(1)$ calculated values do not match the $O(N)$, "true" values, the class throws an exception, which is caught by the CATCH testing framework and causes the test to fail.

In addition, the AVL map maintains a balance factor between -1 and 1 at every subtree. I included another recursive function, enabled when `__DEBUG__==true`, to check for any subtrees that have $\text{abs}(\text{balance factor}) > 1$. If it finds any, it throws an exception which is caught by the CATCH testing framework.

operation_successes.cpp

Within the success operations scenario, I test 4 versions of the map - one for each of the keys to be supported (int, double, string, and c-string). To each instance, I start by filling the map with a bunch of items, clearing it, then filling it up again. The map should then report the correct size. I check that several keys which are expected to exist in the map actually do exist (including the lowest possible key, the highest possible key, and one in the middle).

I check the print() function by routing it to an output string stream and count the number of hyphens in the output, which indicate empty slots. Since $\text{load factor} = \text{occupied buckets} /$

capacity, we can get the number of unoccupied buckets as $\text{capacity} * (1 - \text{load})$. This should equal the number of hyphens in the `print()` output.

I then attempt to `remove()` several keys which are known to exist, and check that their associated values are what were expected. After these items are removed, I try to both `search()` and `remove()` them, which should all return false.

I test `remove_random()` in a similar way to the preceding `remove()` check. I remove a random key, then try to `search()` for it and `remove()` it. Both checks should fail and return a value less than zero.

tree__structure.cpp

Because I was able to successfully implement the pretty-print bonus method, it is fairly straightforward to verify that the underlying AVL map successfully maintains the correct BST tree structure. I print the tree structure to a file, then make a `system()` call to pipe that through a few command line utilities that extract just the numerical key (I say numerical, but that number could be in the form of a string of `const char*`). This results in the natural order of the keys as they exist in the tree. I pipe those through `uniq` and `sort` to remove duplicates and induce expected sorting order. If the tree exhibits the correct structure, the output of those two operations will match verbatim.

I start by filling the map to half capacity, verify that `load()` returns 0.5, then check validate the tree structure. Afterwards, I fill the tree to full capacity and then delete half the nodes via `remove()` and `remove_random()` operations in an attempt to destabilize tree structure. Then I validate the tree structure again.

This series of checks is implemented for each of the four supported key types.

part4_bonus/part4bonus.pdf

part4_bonus/checklist.txt

AVL Tree written by Nickerson, Paul
COP 3530, 2014F 1087

=====
Part IV BONUS: AVL Tree

=====
My MAP implementation uses the data structure described in the part IV
instructions and conforms to the technique required for this map
variety: yes

My MAP implementation 100% correctly implements AVL tree behavior: yes

My MAP implementation 100% correctly supports the following key types:

- * signed int: yes
- * double: yes
- * c-string: yes
- * std::string: yes

My MAP implementation 100% correctly supports the ANY value type: yes

My MAP implementation 100% correctly supports the following methods
as described in part IV:

- * insert: yes
- * remove: yes
- * search: yes
- * search: yes
- * clear: yes
- * is_empty: yes
- * capacity: yes
- * size: yes
- * load: yes
- * print: yes
- * cluster_distribution(): yes
- * remove_random(): yes

My MAP implementation 100% correctly implements the bonus print(): yes

=====
FOR ALL PARTS

=====
My MAP implementation compiles correctly using g++ v4.8.2 on the

OpenBSD VM: yes

My TEST compiles correctly using g++ v4.8.2 on the OpenBSD VM: yes

My TESTS run correctly on the OpenBSD VM: yes

I affirm that all the responses I have provided above are 100% true.
Should it be determined that any are not 100% true, I agree to take a 0
(zero) on the assignment: yes

I affirm that I am the sole author of this AVL Tree
and the associated tests.
Paul Nickerson, Dec 2 in COP3520 section 1087

How to compile and run my unit tests on the OpenBSD VM
cd part4bonus/source
./compile.sh
./run_tests > output.txt

common/common.h

common/common.h

```
1  #ifndef _COMMON_H_
2  #define _COMMON_H_
3
4  #ifndef _DEBUG_
5      //various internal integrity checks that can be expensive, we want to disable
6      //them in production
7      #define _DEBUG_ false
8  #endif
9
10 #include <string.h>
11 #include <limits>
12 #include <stdexcept>
13 #include <ostream>
14 #include <cmath>
15
16 namespace cop3530 {
17     inline double lg(size_t i) {
18         return std::log(i) / std::log(2);
19     }
20     inline size_t rand_i(size_t max) {
21         size_t bucket_size = RAND_MAX / max;
22         size_t num_buckets = RAND_MAX / bucket_size;
23         size_t big_rand;
24         do {
25             big_rand = rand();
26         } while(big_rand >= num_buckets * bucket_size);
27         return big_rand / bucket_size;
28     }
29
30 namespace hash_utils {
31     static constexpr size_t max_size_t = std::numeric_limits<size_t>::max();
32     static constexpr size_t primes[] = { //from algorithms in c++, helps us to
33         choose a prime-number map capacity
34         251,
35         509,
36         1021,
37         2039,
38         4093,
39         8191,
40         16381,
41         32749,
42         65521,
43         131071,
44         262193,
45         524287,
46         1048573,
47         2097143,
```

```

46         4194301,
47         8388593,
48         16777213,
49         33554393,
50         67108859,
51         134217689,
52         268435399,
53         536870909,
54         1073741789,
55         2147483647
56     };
57     struct ClusterInventory {
58         size_t cluster_size;
59         size_t num_instances;
60         struct cluster_size_less_predicate {
61             bool operator()(ClusterInventory const& cluster1, ClusterInventory
62                 const& cluster2) {
63                 return cluster1.cluster_size < cluster2.cluster_size;
64             }
65         };
66     };
67     inline size_t str_to_numeric(const char* str) {
68         unsigned int base = 257; //prime number chosen near an 8-bit character
69         size_t numeric = 0;
70         for (; *str != 0; ++str)
71             numeric = numeric * base + *str;
72         return numeric;
73     }
74     namespace functors {
75
76         struct map_capacity_planner {
77             size_t operator()(size_t min_capacity) {
78                 for (int i = 0; i != 24; ++i)
79                     if (min_capacity < primes[i])
80                         return primes[i];
81                 throw std::domain_error("Provided min capacity too large.
82                     Consider extending the list of prime numbers");
83             }
84         };
85         struct compare {
86             int operator()(const char* a, const char* b) const {
87                 int cmp = strcmp(a, b);
88                 return (cmp < 0 ? -1 :
89                     (cmp > 0 ? 1 : 0));
90             }
91             int operator()(double a, double b) const {
92                 return (a < b ? -1 :
93                     (a > b ? 1 : 0));
94             }
95             int operator()(std::string const& a, std::string const& b) const {
96                 return (a < b ? -1 :

```

```

96             (a > b ? 1 : 0));
97     }
98     int operator()(int a, int b) const {
99         return (a < b ? -1 :
100             (a > b ? 1 : 0));
101     }
102 };
103 namespace primary_hashes {
104     struct hash_basic {
105         //this is such a stupid hash method, but unlike my pathetic attempts
106         //at implementing
107         //various other hashing methods, it works and is generalizable to
108         //all the required key
109         //types. together with double hashing it should make for a passable
110         //hashing routine.
111     public:
112         size_t operator()(const char* key) const {
113             return str_to_numeric(key);
114         }
115         size_t operator()(double key) const {
116             return static_cast<size_t>(std::fmod(key, max_size_t));
117         }
118         size_t operator()(int key) const {
119             return static_cast<size_t>(key);
120         }
121         size_t operator()(std::string const& key) const {
122             const char* c_key = key.c_str();
123             return operator()(c_key);
124         }
125     };
126 }
127 namespace secondary_hashes {
128     struct linear_probe {
129         bool changes_with_probe_attempt() const {
130             return false;
131         }
132     template<typename T>
133     size_t operator()(T unused, size_t probe_attempt) const {
134         return 1;
135     }
136 };
137 struct quadratic_probe {
138     bool changes_with_probe_attempt() const {
139         return true;
140     }
141     template<typename T>
142     size_t operator()(T unused, size_t probe_attempt) const {
143         return probe_attempt;
144     }
145 };
146 struct hash_double {
147 private:

```

```

145         size_t hash_numeric(size_t numeric) const {
146             size_t hash = numeric % 97; //simple modulus using a prime
                number (from algorithms in c++)
147             //the second hash may not be zero (will cause an infinite
                loop).
148             //also, hash must be relatively prime to map_capacity so that
                every slot can be hit.
149             //map capacity is a prime number based chosen from the table,
                so any value less than
150             //map capacity should work
151             return hash;
152         }
153     public:
154         bool changes_with_probe_attempt() const {
155             return false;
156         }
157         size_t operator()(const char* key, size_t) const {
158             size_t numeric = str_to_numeric(key);
159             return hash_numeric(numeric);
160         }
161         size_t operator()(double key, size_t) const {
162             return hash_numeric(key);
163         }
164         size_t operator()(int key, size_t) const {
165             return hash_numeric(key);
166         }
167         size_t operator()(std::string key, size_t) const {
168             const char* c_key = key.c_str();
169             return operator()(c_key, 0);
170         }
171     };
172 }
173
174
175 template<typename T>
176 class GenericContainer {
177     /*
178     for the types we need to support other than const char* (ie int,
        double, and std::string),
179     we can pass these around willy-nilly. for const char*, handled
        below, we will obtain our
180     own copy of the character array by wrapping it in a std::string
181     */
182 private:
183     T raw;
184     functors::compare compare;
185 public:
186     GenericContainer(const T& val): raw(val) {}
187     GenericContainer() = default;
188     GenericContainer& operator=(GenericContainer const& rhs) = delete;
189     T operator()() const {
190         return raw;

```

```

191     }
192     T copy() const {
193         return raw;
194     }
195     void reset(const T& val) {
196         raw = val;
197     }
198     int compare_to(GenericContainer const& other) const {
199         return compare(raw, other.raw);
200     }
201 };
202 template<>
203 class GenericContainer<const char*> {
204     /*
205         class template specialization for character arrays, stores a local
206         copy of the character array
207     */
208 private:
209     char* raw = nullptr;
210     functors::compare compare;
211 public:
212     GenericContainer(const char* val) {
213         reset(val);
214     }
215     GenericContainer() = default;
216     const char* operator>()() const {
217         return raw;
218     }
219     const char* copy() const {
220         if (raw == nullptr) return nullptr;
221         size_t len = strlen(raw);
222         char* new_str = new char[len + 1];
223         strncpy(new_str, raw, len);
224         new_str[len] = 0;
225         return new_str;
226     }
227     void reset(const char* val) {
228         if (raw) {
229             delete raw;
230             raw = nullptr;
231         }
232         if (val != nullptr) {
233             size_t len = strlen(val);
234             raw = new char[len + 1];
235             strncpy(raw, val, len);
236             raw[len] = 0;
237         }
238     }
239     int compare_to(GenericContainer const& other) const {
240         return compare(raw, other.raw);
241     }
242 };

```

```

242
243     template<typename key_type,
244             typename primary_hash =
                hash_utils::functors::primary_hashes::hash_basic,
245             typename secondary_hash =
                hash_utils::functors::secondary_hashes::hash_double>
246     class Key {
247     private:
248         GenericContainer<key_type> raw_key;
249         primary_hash hasher1;
250         secondary_hash hasher2;
251         size_t hash1_val;
252         size_t hash2_val;
253     public:
254         Key& operator=(Key const& rhs) {
255             if (&rhs == this)
256                 return *this;
257             reset(rhs.raw_key());
258         }
259         bool operator==(Key const& rhs) const {
260             return raw_key.compare_to(rhs.raw_key) == 0;
261         }
262         bool operator<(Key const& rhs) const {
263             return raw_key.compare_to(rhs.raw_key) == -1;
264         }
265         bool operator>(Key const& rhs) const {
266             return raw_key.compare_to(rhs.raw_key) == 1;
267         }
268         bool operator!=(Key const& rhs) const {
269             return ! operator==(rhs);
270         }
271         size_t hash(size_t map_capacity, size_t probe_attempt) const {
272             size_t local_hash2_val;
273             if (probe_attempt != 0 && hasher2.changes_with_probe_attempt())
274             {
275                 //if the hashing function value is dependent on the probe attempt
276                 //(eg quadratic probing), then we need to retrieve the new value*/
277                 local_hash2_val = hasher2(raw_key(), probe_attempt);
278             } else {
279                 //otherwise we can just use the value we have stored
280                 local_hash2_val = hash2_val;
281             }
282             return (hash1_val + probe_attempt * local_hash2_val) % map_capacity;
283         }
284         key_type raw() const {
285             return raw_key();
286         }
287         key_type raw_copy() const {
288             //this is what is returned to the client, who is responsible for
289             deleting it if its, eg a pointer to a character array
290             return raw_key.copy();

```



```

291     template<typename T>
292     void reset(T key) {
293         raw_key.reset(key);
294         size_t base_probe_attempt = 0;
295         hash1_val = hasher1(key);
296         hash2_val = hasher2(key, base_probe_attempt);
297     }
298     void reset(const char* key) {
299         raw_key.reset(key);
300         if (key != nullptr) {
301             size_t base_probe_attempt = 0;
302             hash1_val = hasher1(key);
303             hash2_val = hasher2(key, base_probe_attempt);
304         }
305     }
306     explicit Key(key_type const& key): raw_key(key) {
307         reset(key);
308     }
309     Key() = default;
310 };
311 template <typename value_type>
312 class Value {
313 private:
314     functors::compare compare;
315     GenericContainer<value_type> raw_value;
316 public:
317     Value& operator=(Value const& rhs) {
318         if (&rhs == this)
319             return *this;
320         reset(rhs.raw_value());
321     }
322     bool operator==(Value const& rhs) const {
323         return compare(raw_value(), rhs.raw_value());
324     }
325     bool operator==(value_type const& rhs) const {
326         return compare(raw_value(), rhs) == 0;
327     }
328     value_type raw() const {
329         return raw_value();
330     }
331     value_type raw_copy() const {
332         //this is what is returned to the client, who is responsible for
333         //deleting it if its, eg a pointer to a character array
334         return raw_value.copy();
335     }
336     void reset(value_type value) {
337         raw_value.reset(value);
338     }
339     explicit Value(value_type const& value): raw_value(value) {}
340     Value() = default;
341 };

```

```
342 }  
343  
344 #endif
```

priority_queue.h

priority_queue.h

```
1  #ifndef _PRIORITY_QUEUE_H_
2  #define _PRIORITY_QUEUE_H_
3
4  #include "SDAL.h"
5  #include "common.h"
6
7  namespace cop3530 {
8      //this class allows efficient sorting clusters by size for the
9      //cluster_distribution functions
10     template<typename T,
11             typename PriorityCompare =
12             cop3530::hash_utils::ClusterInventory::cluster_size_less_predicate>
13     class priority_queue {
14     private:
15         PriorityCompare first_arg_higher_priority;
16         //SDAL has all the benefits of std::vector (ie fast random access and
17         //automatic resizing)
18         //while having the added benefit of being legal to use in cop3530
19         SDAL<T> tree;
20         size_t num_items = 0;
21         void fix_up(size_t index) {
22             while (index > 1
23                 && first_arg_higher_priority(tree[index], tree[index / 2]))
24             {
25                 std::swap(tree[index / 2], tree[index]);
26                 index /= 2;
27             }
28         }
29         void fix_down() {
30             size_t parent_index = 1;
31             while (2 * parent_index <= num_items) {
32                 size_t left_index = 2 * parent_index;
33                 size_t right_index = left_index + 1;
34                 size_t higher_priority_index = left_index;
35                 if (right_index <= num_items
36                     && first_arg_higher_priority(tree[right_index], tree[left_index]))
37                 {
38                     higher_priority_index = right_index;
39                 }
40                 if ( ! first_arg_higher_priority(tree[higher_priority_index],
41                     tree[parent_index]))
42                     //no more items to elevate
43                     break;
44                 std::swap(tree[parent_index], tree[higher_priority_index]);
45                 parent_index = higher_priority_index;
46             }
47         }
48     }
```

```

44     public:
45         priority_queue(size_t preallocation_size = 100): tree(preallocation_size +
46             1) {
47             T empty_item;
48             tree.push_back(empty_item);
49         }
50         priority_queue(priority_queue const& src) {
51             tree = src.tree;
52             num_items = src.num_items;
53         }
54         T get_next_item() {
55             std::swap(tree[1], tree[num_items]);
56             T ret = tree[num_items--];
57             fix_down();
58             return ret;
59         }
60         void add_to_queue(T const& item) {
61             tree.push_back(item);
62             num_items++;
63             fix_up(num_items);
64         }
65         size_t size() {
66             return num_items;
67         }
68         bool empty() {
69             return num_items == 0;
70         }
71     };
72
73 #endif // _PRIORITY_QUEUE_H_

```

common/unit_test_utils.h

common/unit_test_utils.h

```
1  #ifndef _UNIT_TEST_UTILS_H_
2  #define _UNIT_TEST_UTILS_H_
3
4  #include <iostream>
5  #include <string>
6  #include <fstream>
7
8  namespace cop3530 {
9      namespace unit_test_utils {
10         inline std::string guid() {
11             std::string ret = "";
12             for (size_t i = 0; i != 32; ++i) {
13                 size_t rnd = cop3530::rand_i(16);
14                 if (rnd < 10)
15                     ret += std::string(1, '0' + rnd);
16                 else
17                     ret += std::string(1, 'A' + rnd - 10);
18             }
19             return ret;
20         }
21         inline std::string get_tmp_filename() {
22             return std::string("/tmp/") + guid() + std::string(".out");
23         }
24         inline void delete_file(std::string file_path) {
25             system((std::string("rm ") + file_path + std::string(" 2>&1 >>
26                 /tmp/debug")).c_str());
27         }
28         inline std::string shell_cmd(std::string cmd) {
29             std::string shell_script_file = get_tmp_filename();
30             std::string output_file = get_tmp_filename();
31             std::ofstream shell_script_out(shell_script_file);
32             shell_script_out << "#!/bin/sh" << std::endl << cmd << std::endl;
33             shell_script_out.close();
34             std::string chmod_cmd = std::string("chmod +x ") + shell_script_file;
35             system(chmod_cmd.c_str());
36             std::string invoke_cmd = shell_script_file + std::string(" > ") +
37                 output_file;
38             system(invoke_cmd.c_str());
39             std::ifstream read_output(output_file);
40             std::ostringstream oss;
41             std::string tmp;
42             while (std::getline(read_output, tmp)) {
43                 oss << tmp;
44                 if ( ! read_output.eof())
45                     oss << "\n";
46             }
47             delete_file(output_file);
48         }
49     }
50 }
```

```
46         delete_file(shell_script_file);
47         return oss.str();
48     }
49 }
50 }
51
52 #endif // _UNIT_TEST_UTILS_H_
```

part4_bonus/source/bst.h

part4_bonus/source/bst.h

```
1  #ifndef _BST_H_
2  #define _BST_H_
3
4  #include <cstdlib>
5  #include <sstream>
6  #include "../common/CDAL.h"
7  #include "../common/common.h"
8  #include "../common/priority_queue.h"
9
10 namespace cop3530 {
11     template<typename key_type,
12             typename value_type>
13     class BST {
14     protected: //let RBST and AVL inherit everything
15         typedef hash_utils::ClusterInventory ClusterInventory;
16         typedef hash_utils::Key<key_type> Key;
17         typedef hash_utils::Value<value_type> Value;
18         struct Node;
19         typedef Node* link;
20         struct Node {
21             Key key;
22             Value value;
23             size_t num_children;
24             size_t left_index;
25             size_t right_index;
26             size_t height; //height tracking coded in this class, but not used (for
                             //AVL, which is this class with self-balancing)
27             bool is_occupied;
28             size_t validate_children_count_recursive(Node* nodes) {
29                 //this function is for debugging purposes, does recursive traversal
                 //to find the correct number of children
30                 size_t child_count = 0;
31                 if (left_index)
32                     child_count += 1 +
                        nodes[left_index].validate_children_count_recursive(nodes);
33                 if (right_index)
34                     child_count += 1 +
                        nodes[right_index].validate_children_count_recursive(nodes);
35                 if (child_count != num_children) {
36                     std::ostringstream msg;
37                     msg << "Manually counted children, " << child_count << ",
                        different than child count, " << num_children;
38                     throw std::logic_error(msg.str());
39                 }
40                 return child_count;
41             }
42             size_t get_height_recursive(Node* nodes) {
```

```

43         //this function is for debugging purposes, does recursive traversal
           to find the correct height
44         size_t left_height = 0, right_height = 0;
45         size_t calculated_height = 0;
46         if (left_index)
47             left_height = nodes[left_index].get_height_recursive(nodes);
48         if (right_index)
49             right_height = nodes[right_index].get_height_recursive(nodes);
50         calculated_height = 1 + std::max(left_height, right_height);
51         return calculated_height;
52     }
53     void update_height(Node* nodes) {
54         //note: this method depends on the left and right subtree heights
           being correct
55         size_t left_height = 0, right_height = 0;
56         if (left_index)
57             left_height = nodes[left_index].height;
58         if (right_index)
59             right_height = nodes[right_index].height;
60         height = 1 + std::max(left_height, right_height);
61         if (_DEBUG_) {
62             size_t calculated_height = get_height_recursive(nodes);
63             if (calculated_height != height) {
64                 std::ostringstream msg;
65                 msg << "Manually calculated height, " << calculated_height <<
66                     ", different than tracked height, " << height;
67                 throw std::logic_error(msg.str());
68             }
69         }
70     void disable_and_adopt_free_tree(size_t free_index) {
71         is_occupied = false;
72         height = 0;
73         num_children = 0;
74         right_index = 0;
75         left_index = free_index;
76     }
77     void reset_and_enable(Key const& new_key, Value const& new_value) {
78         is_occupied = true;
79         height = 1; //self
80         left_index = right_index = 0;
81         num_children = 0;
82         key = new_key;
83         value = new_value;
84     }
85     int balance_factor(const Node* nodes) const {
86         size_t left_height = 0, right_height = 0;
87         if (left_index)
88             left_height = nodes[left_index].height;
89         if (right_index)
90             right_height = nodes[right_index].height;

```



```

91         return static_cast<long int>(left_height) - static_cast<long
92             int>(right_height);
93     }
94 };
95 Node* nodes; /***note: array is 1-based so leaf nodes have child indices
96     set to zero
97 size_t free_index;
98 size_t root_index;
99 size_t curr_capacity;
100 virtual size_t remove_smallest_key_node_index(size_t& subtree_root_index) {
101     //returns the index of the node with the smallest key, while
102     //setting its parent's left child index to the smallest key node's
103     //right child index. recursion downward through this function updates
104     //the heights of the nodes it traverses
105     Node& subtree_root = nodes[subtree_root_index];
106     size_t smallest_key_node_index = 0;
107     if (subtree_root_index == 0) {
108         throw std::logic_error("Expected to find a valid node, but didn't");
109     } else {
110         if (subtree_root.left_index) {
111             smallest_key_node_index =
112                 remove_smallest_key_node_index(subtree_root.left_index);
113             subtree_root.num_children--;
114             subtree_root.update_height(nodes);
115         } else {
116             smallest_key_node_index = subtree_root_index;
117             subtree_root_index = subtree_root.right_index;
118         }
119     }
120     return smallest_key_node_index;
121 }
122 virtual size_t remove_largest_key_node_index(size_t& subtree_root_index) {
123     //returns the index of the node with the largest key, while
124     //setting its parent's right child index to the largest key node's
125     //left child index. recursion downward through this function updates
126     //the heights of the nodes it traverses
127     Node& subtree_root = nodes[subtree_root_index];
128     size_t largest_key_node_index = 0;
129     if (subtree_root_index == 0) {
130         throw std::logic_error("Expected to find a valid node, but didn't");
131     } else {
132         if (subtree_root.right_index) {
133             largest_key_node_index =
134                 remove_largest_key_node_index(subtree_root.right_index);
135             subtree_root.num_children--;
136             subtree_root.update_height(nodes);
137         } else {
138             largest_key_node_index = subtree_root_index;
139             subtree_root_index = subtree_root.left_index;
140         }
141     }
142     return largest_key_node_index;

```

```

139     }
140     virtual void remove_node(size_t& subtree_root_index) {
141         Node& subtree_root = nodes[subtree_root_index];
142         size_t index_to_delete = subtree_root_index;
143         if (subtree_root.right_index || subtree_root.left_index) {
144             //subtree has at least one child
145             if (subtree_root.right_index)
146                 //replace the root with the smallest-keyed node in the right
147                 subtree
148                 subtree_root_index =
149                     remove_smallest_key_node_index(subtree_root.right_index);
150             else if (subtree_root.left_index)
151                 //replace the root with the largest-keyed node in the left subtree
152                 subtree_root_index =
153                     remove_largest_key_node_index(subtree_root.left_index);
154             //have the new root adopt the old root's children
155             Node& new_root = nodes[subtree_root_index];
156             new_root.left_index = subtree_root.left_index;
157             new_root.right_index = subtree_root.right_index;
158             //the new root has the same number of children as the old root,
159             //minus one
160             new_root.num_children = subtree_root.num_children - 1;
161             //removing the smallest/largest-keyed node from the old root has the
162             //effect of
163             //updating the heights of the old root's relevant subtrees (which
164             //the new root
165             //just adopted), so we can update the new root's height now
166             new_root.update_height(nodes);
167         } else
168             //neither subtree exists, so just delete the node
169             subtree_root_index = 0;
170         //node has been disowned by all ancestors, and has disowned all
171         //descendents, so free it
172         add_node_to_free_tree(index_to_delete);
173     }
174     virtual int do_remove(size_t nodes_visited, //starts at 0 when this
175                           function is first called (ie does not include current node visitation)
176                           size_t& subtree_root_index,
177                           Key const& key,
178                           Value& value,
179                           bool& found_key)
180     {
181         if (subtree_root_index != 0) {
182             Node& subtree_root = nodes[subtree_root_index];
183             ++nodes_visited;
184             //keep going down to the base of the tree
185             if (key < subtree_root.key) {
186                 nodes_visited = do_remove(nodes_visited, subtree_root.left_index,
187                                           key, value, found_key);
188             }
189             if (found_key) {
190                 //found the desired node and delete it
191                 subtree_root.num_children--;
192             }
193         }
194     }

```

```

182         //left child changed, so recompute subtree height
183         subtree_root.update_height(nodes);
184     }
185     } else if (key > subtree_root.key) {
186         nodes_visited = do_remove(nodes_visited,
187             subtree_root.right_index, key, value, found_key);
188         if (found_key) {
189             //found the desired node and delete it
190             subtree_root.num_children--;
191             //right child changed, so recompute subtree height
192             subtree_root.update_height(nodes);
193         }
194     } else if (key == subtree_root.key) {
195         //found key, remove the node
196         found_key = true;
197         value = subtree_root.value;
198         remove_node(subtree_root_index);
199     } else {
200         throw std::logic_error("Unexpected compare result");
201     }
202 }
203 return nodes_visited;
204 }
205 void write_subtree_buffer(size_t subtree_root_index,
206     CDAL<std::string>& buffer_lines,
207     size_t root_line_index,
208     size_t lbound_line_index /*inclusive*/,
209     size_t ubound_line_index /*exclusive*/) const
210 {
211     Node subtree_root = nodes[subtree_root_index];
212     std::ostringstream oss;
213     //print the node
214     //todo: fix this to only print the key
215     oss << "[" << subtree_root.key.raw() << "]";
216     buffer_lines[root_line_index] += oss.str();
217     //print the right descendents
218     if (subtree_root.right_index > 0) {
219         //at least 1 right child
220         size_t top_dashes = 1;
221         Node const& right_child = nodes[subtree_root.right_index];
222         if (right_child.left_index > 0) {
223             //right child has at least 1 left child
224             Node const& right_left_child = nodes[right_child.left_index];
225             top_dashes += 2 * (1 + right_left_child.num_children);
226         }
227         size_t top_line_index = root_line_index - 1;
228         while (top_line_index >= root_line_index - top_dashes)
229             buffer_lines[top_line_index--] += "| ";
230         size_t right_child_line_index = top_line_index;
231         buffer_lines[top_line_index--] += "+--";
232         while (top_line_index >= lbound_line_index)
233             buffer_lines[top_line_index--] += " ";

```

```

233         write_subtree_buffer(subtree_root.right_index,
234                               buffer_lines,
235                               right_child_line_index,
236                               lbound_line_index,
237                               root_line_index);
238     }
239     //print the left descendents
240     if (subtree_root.left_index > 0) {
241         //at least 1 left child
242         size_t bottom_dashes = 1;
243         Node const& left_child = nodes[subtree_root.left_index];
244         if (left_child.right_index > 0) {
245             //left child has at least 1 right child
246             Node const& left_right_child = nodes[left_child.right_index];
247             bottom_dashes += 2 * (1 + left_right_child.num_children);
248         }
249         size_t bottom_line_index = root_line_index + 1;
250         while (bottom_line_index <= root_line_index + bottom_dashes)
251             buffer_lines[bottom_line_index++] += "| ";
252         size_t left_child_line_index = bottom_line_index;
253         buffer_lines[bottom_line_index++] += "+--";
254         while (bottom_line_index < ubound_line_index)
255             buffer_lines[bottom_line_index++] += " ";
256         write_subtree_buffer(subtree_root.left_index,
257                               buffer_lines,
258                               left_child_line_index,
259                               root_line_index + 1,
260                               ubound_line_index);
261     }
262 }
263 void add_node_to_free_tree(size_t node_index) {
264     nodes[node_index].disable_and_adopt_free_tree(free_index);
265     free_index = node_index;
266 }
267 size_t procure_node(Key const& key, Value const& value) {
268     //updates the free index to the first free node's left child (while
269     //transforming that first free
270     //node to an enabled node with the specified key/value) and returns the
271     //index of what was the last
272     //free index
273     size_t node_index = free_index;
274     free_index = nodes[free_index].left_index;
275     Node& n = nodes[node_index];
276     n.reset_and_enable(key, value);
277     return node_index;
278 }
279 virtual int insert_at_leaf(size_t nodes_visited, //starts at 0 when this
280                           //function is first called (ie does not include current node visitation)
281                           size_t& subtree_root_index,
282                           Key const& key,
283                           Value const& value,
284                           bool& found_key)

```

```

282     {
283         if (subtree_root_index == 0) {
284             //key not found
285             subtree_root_index = procure_node(key, value);
286         } else {
287             //parent was not a leaf
288             //keep going down to the base of the tree
289             Node& subtree_root = nodes[subtree_root_index];
290             ++nodes_visited;
291             if (key < subtree_root.key) {
292                 nodes_visited = insert_at_leaf(nodes_visited,
293                     subtree_root.left_index, key, value, found_key);
294                 if ( ! found_key) {
295                     //given key is unique to the tree, so a new node was added
296                     subtree_root.num_children++;
297                     subtree_root.update_height(nodes);
298                 }
299             } else if (key > subtree_root.key) {
300                 nodes_visited = insert_at_leaf(nodes_visited,
301                     subtree_root.right_index, key, value, found_key);
302                 if ( ! found_key) {
303                     //given key is unique to the tree, so a new node was added
304                     subtree_root.num_children++;
305                     subtree_root.update_height(nodes);
306                 }
307             } else if (key == subtree_root.key) {
308                 //found key, replace the value
309                 subtree_root.value = value;
310                 found_key = true;
311             } else {
312                 throw std::logic_error("Unexpected compare result");
313             }
314         }
315         return nodes_visited;
316     }
317 }
318
319 void rotate_left(size_t& subtree_root_index) {
320     Node& subtree_root = nodes[subtree_root_index];
321     size_t right_child_index = subtree_root.right_index;
322     Node& right_child = nodes[right_child_index];
323
324     //original root adopts the right child's left subtree
325     subtree_root.right_index = right_child.left_index;
326     //original root adopted a subtree (whose height did not change), so
327     //update its height
328     subtree_root.update_height(nodes);
329
330     //right child adopts original root and its children
331     right_child.left_index = subtree_root_index;
332     //right child (new root) adopted the original root (whose height has
333     //been updated), so update its height
334     right_child.update_height(nodes);

```

```

329         //since right child took the subtree root's place, it has the same
           number of children as the original root
330         right_child.num_children = subtree_root.num_children;
331
332         //root has new children, so update that counter (done after changing the
           right child's children counter
333         //because that depends on the original root's counter)
334         subtree_root.num_children = 0;
335         if (subtree_root.left_index != 0)
336             subtree_root.num_children += 1 +
           nodes[subtree_root.left_index].num_children;
337         if (subtree_root.right_index != 0)
338             subtree_root.num_children += 1 +
           nodes[subtree_root.right_index].num_children;
339
340         //set the right child as the new root
341         subtree_root_index = right_child_index;
342     }
343     void rotate_right(size_t& subtree_root_index) {
344         Node& subtree_root = nodes[subtree_root_index];
345         size_t left_child_index = subtree_root.left_index;
346         Node& left_child = nodes[left_child_index];
347
348         //original root adopts the left child's right subtree
349         subtree_root.left_index = left_child.right_index;
350         //original root adopted a subtree (whose height did not change), so
           update its height
351         subtree_root.update_height(nodes);
352
353         //left child adopts original root and its children
354         left_child.right_index = subtree_root_index;
355         //left child (new root) adopted the original root (whose height has been
           updated), so update its height
356         left_child.update_height(nodes);
357         //since left child took the subtree root's place, it has the same number
           of children as the original root
358         left_child.num_children = subtree_root.num_children;
359
360         //root has new children, so update that counter (done after changing the
           left child's children counter
361         //because that depends on the original root's counter)
362         subtree_root.num_children = 0;
363         if (subtree_root.left_index != 0)
364             subtree_root.num_children += 1 +
           nodes[subtree_root.left_index].num_children;
365         if (subtree_root.right_index != 0)
366             subtree_root.num_children += 1 +
           nodes[subtree_root.right_index].num_children;
367
368         //set the left child as the new root
369         subtree_root_index = left_child_index;
370     }

```

```

371 int do_search(size_t nodes_visited, //starts at 0 when this function is
    first called (ie does not include current node visitation)
372             size_t subtree_root_index,
373             Key const& key,
374             Value& value,
375             bool& found_key) const
376 {
377     if (subtree_root_index != 0) {
378         Node const& subtree_root = nodes[subtree_root_index];
379         ++nodes_visited;
380         if (key < subtree_root.key) {
381             nodes_visited = do_search(nodes_visited, subtree_root.left_index,
    key, value, found_key);
382         } else if (key > subtree_root.key) {
383             nodes_visited = do_search(nodes_visited,
    subtree_root.right_index, key, value, found_key);
384         } else if (key == subtree_root.key) {
385             //found key, replace the value
386             value = subtree_root.value;
387             found_key = true;
388         } else {
389             throw std::logic_error("Unexpected compare result");
390         }
391     }
392     return nodes_visited;
393 }
394 void prepare_cluster_distribution(size_t subtree_root_index,
    size_t curr_height, //includes the height of
    the current node, ie assumes current node
    exists
    size_t cluster_counter[])
396 {
397     Node const& subtree_root = nodes[subtree_root_index];
398     if ( ! subtree_root.left_index && ! subtree_root.right_index)
399         //at a leaf node
400         cluster_counter[curr_height]++;
401     else {
402         if (subtree_root.left_index)
403             prepare_cluster_distribution(subtree_root.left_index, curr_height
    + 1, cluster_counter);
404         if (subtree_root.right_index)
405             prepare_cluster_distribution(subtree_root.right_index,
    curr_height + 1, cluster_counter);
406     }
407 }
408
409 void remove_ith_node_inorder(size_t& subtree_root_index,
    size_t& ith_node_to_delete,
    Key& key)
410 {
411     Node& subtree_root = nodes[subtree_root_index];
412     if (subtree_root.left_index)
413

```

```

416         remove_ith_node_inorder(subtree_root.left_index, ith_node_to_delete,
417                                 key);
418     if (ith_node_to_delete == 0)
419         //deleted node in child subtree; nothing more to do
420         return;
421     if (--ith_node_to_delete == 0) {
422         //delete the current node
423         value_type dummy_val;
424         remove(subtree_root.key.raw_copy(), dummy_val);
425         key = subtree_root.key;
426         return;
427     }
428     if (subtree_root.right_index)
429         remove_ith_node_inorder(subtree_root.right_index,
430                                 ith_node_to_delete, key);
431 }
432
433 public:
434     /*
435     The constructor will allocate an array of capacity (binary
436     tree) nodes. Then make a chain from all the nodes (e.g.,
437     make node 2 the left child of node 1, make node 3 the left
438     child of node 2, &c. this is the initial free list.
439     */
440     BST(size_t capacity):
441         curr_capacity(capacity)
442     {
443         if (capacity == 0) {
444             throw std::domain_error("capacity must be at least 1");
445         }
446         nodes = new Node[capacity + 1];
447         clear();
448     }
449     /*
450     if there is space available, adds the specified key/value-pair to the
451     tree
452     and returns the number of nodes visited, V; otherwise returns -1 * V. If
453     an
454     item already exists in the tree with the same key, replace its value.
455     */
456     virtual int insert(key_type const& key, value_type const& value) {
457         if (size() == capacity())
458             //no more space
459             return -1 * size();
460         bool found_key = false;
461         Key k(key);
462         Value v(value);
463         int nodes_visited = insert_at_leaf(0, root_index, k, v, found_key);
464         if (_DEBUG_)
465             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
466         return nodes_visited;
467     }

```



```

464     /*
465         if there is an item matching key, removes the key/value-pair from the
            tree, stores
466         it's value in value, and returns the number of probes required, V;
            otherwise returns -1 * V.
467     */
468     virtual int remove(key_type const& key, value_type& value) {
469         if (is_empty())
470             return 0;
471         bool found_key = false;
472         Key k(key);
473         Value v(value);
474         int nodes_visited = do_remove(0, root_index, k, v, found_key);
475         if (_DEBUG_)
476             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
477         if (found_key)
478             value = v.raw_copy();
479         return found_key ? nodes_visited : -1 * nodes_visited;
480     }
481     /*
482         if there is an item matching key, stores it's value in value, and
            returns the number
483         of nodes visited, V; otherwise returns -1 * V. Regardless, the item
            remains in the tree.
484     */
485     virtual int search(key_type const& key, value_type& value) {
486         if (is_empty())
487             return 0;
488         bool found_key = false;
489         Key k(key);
490         Value v(value);
491         int nodes_visited = do_search(0, root_index, k, v, found_key);
492         if (found_key)
493             value = v.raw_copy();
494         return found_key ? nodes_visited : -1 * nodes_visited;
495     }
496     /*
497         removes all items from the map
498     */
499     virtual void clear() {
500         //Since I use size_t to hold the node indices, I make the node array
501         //1-based, with child index of 0 indicating that the current node is a
            leaf
502         for (size_t i = 1; i != capacity(); ++i)
503             nodes[i].disable_and_adopt_free_tree(i + 1);
504         free_index = 1;
505         root_index = 0;
506     }
507     /*
508         returns true IFF the map contains no elements.
509     */
510     virtual bool is_empty() const {

```

```

511         return size() == 0;
512     }
513     /*
514     returns the number of slots in the backing array.
515     */
516     virtual size_t capacity() const {
517         return curr_capacity;
518     }
519     /*
520     returns the number of items actually stored in the tree.
521     */
522     virtual size_t size() const {
523         if (root_index == 0) return 0;
524         Node const& root = nodes[root_index];
525         return 1 + root.num_children;
526     }
527     /*
528     [not a regular BST operation, but specific to this implementation]
529     returns the tree's load factor: load = size / capacity.
530     */
531     virtual double load() const {
532         return static_cast<double>(size()) / capacity();
533     }
534     /*
535     prints the tree in the following format:
536     +--[tiger]
537     | |
538     | | +--[panther]
539     | | |
540     | +--[ocelot]
541     | |
542     | +--[lion]
543     |
544     [leopard]
545     |
546     | +--[house cat]
547     | |
548     | +--[cougar]
549     | |
550     +--[cheetah]
551     |
552     +--[bobcat]
553     */
554     virtual std::ostream& print(std::ostream& out) const {
555         if (is_empty())
556             return out;
557         size_t num_lines = size() * 2 - 1;
558         //use CDAL here so we can print really super-huge trees where the write
559         //buffer doesn't fit in memory
560         CDAL<std::string> buffer_lines(100000);
561         for(size_t i = 0; i <= num_lines; ++i)
562             buffer_lines.push_back("");

```

```

562     Node const& root = nodes[root_index];
563     size_t root_line_index = 1;
564     if (root.right_index) {
565         root_line_index += 2 * (1 + nodes[root.right_index].num_children);
566     }
567     write_subtree_buffer(root_index, buffer_lines, root_line_index, 1,
568                          num_lines + 1);
569     for (size_t i = 1; i <= num_lines; ++i)
570         out << buffer_lines[i] << std::endl;
571     return out;
572 }
573
574 /*
575  returns a list indicating the number of leaf nodes at each height (since
576  the RBST doesn't exhibit
577  true clustering, but can have degenerate branches).
578 */
579 virtual priority_queue<hash_utils::ClusterInventory> cluster_distribution()
580 {
581     //use an array to count cluster instances, then feed those to a priority
582     //queue and return it.
583     priority_queue<ClusterInventory> cluster_pq;
584     if (is_empty()) return cluster_pq;
585     size_t max_height = nodes[root_index].height;
586     size_t cluster_counter[max_height + 1];
587     for (size_t i = 0; i <= max_height; ++i)
588         cluster_counter[i] = 0;
589     prepare_cluster_distribution(root_index, 1, cluster_counter);
590     for (size_t i = 1; i <= max_height; ++i)
591         if (cluster_counter[i] > 0) {
592             ClusterInventory cluster{i, cluster_counter[i]};
593             cluster_pq.add_to_queue(cluster);
594         }
595     return cluster_pq;
596 }
597
598 /*
599  generate a random number, R, (1,size), and starting with the root (node
600  1), do an in-order
601  traversal to find the R-th occupied node; remove that node (adjusting
602  its children accordingly),
603  and return its key.
604
605  ***XXX: this likely contains a bug when using const char* keys in that
606  we'll be returning a dangling pointer!!!**
607 */
608 virtual key_type remove_random() {
609     if (size() == 0) throw std::logic_error("Cant remove from an empty map");
610     size_t ith_node_to_delete = 1 + rand_i(size());
611     Key key;
612     remove_ith_node_inorder(root_index, ith_node_to_delete, key);
613     key_type ret = key.raw_copy();

```

```
607         return ret;
608     }
609 };
610 }
611
612 #endif
```

part4_bonus/source/avl.h

part4_bonus/source/avl.h

```
1  #ifndef _AVL_H_
2  #define _AVL_H_
3
4  #include <cstdlib>
5  #include <sstream>
6  #include "../common/CDAL.h"
7  #include "../common/common.h"
8  #include "../common/priority_queue.h"
9  #include "../part4/source/bst.h"
10
11 namespace cop3530 {
12     template<typename key_type,
13             typename value_type>
14     class AVL: public BST<key_type, value_type> {
15     /*
16         The trick to AVL is to perform standard BST operations, but wrap recursive
17         methods that might unbalance
18         the tree with methods that rebalance the tree after performing those
19         operations. Thus the balance factor
20         of any given node stays within [-1, 1]. To that end we simply inherit from
21         a BST base class that tracks
22         changes in subtree height and overwrite the needed virtual methods.
23     */
24     private:
25         using super = BST<key_type, value_type>;
26         using typename super::Node;
27         typedef hash_utils::Key<key_type> Key;
28         typedef hash_utils::Value<value_type> Value;
29         int insert_at_leaf(size_t nodes_visited,
30                           size_t& subtree_root_index,
31                           Key const& key,
32                           Value const& value,
33                           bool& found_key)
34         {
35             nodes_visited = super::insert_at_leaf(nodes_visited, subtree_root_index,
36             key, value, found_key);
37             balance(subtree_root_index);
38             return nodes_visited;
39         }
40         size_t remove_smallest_key_node_index(size_t& subtree_root_index) {
41             size_t smallest_key_node_index =
42                 super::remove_smallest_key_node_index(subtree_root_index);
43             balance(subtree_root_index);
44             return smallest_key_node_index;
45         }
46         size_t remove_largest_key_node_index(size_t& subtree_root_index) {
```

```

42         size_t largest_key_node_index =
43             super::remove_largest_key_node_index(subtree_root_index);
44         balance(subtree_root_index);
45         return largest_key_node_index;
46     }
47     int do_remove(size_t nodes_visited, //starts at 0 when this function is
48         //first called (ie does not include current node visitation)
49         size_t& subtree_root_index,
50         Key const& key,
51         Value& value,
52         bool& found_key)
53     {
54         nodes_visited = super::do_remove(nodes_visited, subtree_root_index, key,
55             value, found_key);
56         balance(subtree_root_index);
57         return nodes_visited;
58     }
59     void balance(size_t& subtree_root_index) {
60         if (subtree_root_index == 0) return;
61         Node& root = this->nodes[subtree_root_index];
62         int root_bal_fact = root.balance_factor(this->nodes);
63         if (root_bal_fact == -2) {
64             //right subtree is too heavy
65             size_t& right_index = root.right_index;
66             Node& right_child = this->nodes[right_index];
67             switch(right_child.balance_factor(this->nodes)) {
68                 case 1:
69                     //right left
70                     this->rotate_right(right_index);
71                     this->rotate_left(subtree_root_index);
72                     break;
73                 case -1:
74                 case 0:
75                     //right right
76                     this->rotate_left(subtree_root_index);
77                     break;
78                 default:
79                     throw std::domain_error(std::string("Unexpected balance factor
80                         with heavy right subtree: ")
81                         +
82                         std::to_string(right_child.balance_factor(this->nodes)));
83             }
84         } else if (root_bal_fact == 2) {
85             //left subtree is too heavy
86             size_t& left_index = root.left_index;
87             Node& left_child = this->nodes[left_index];
88             switch(left_child.balance_factor(this->nodes)) {
89                 case -1:
90                     //left right
91                     this->rotate_left(left_index);
92                     this->rotate_right(subtree_root_index);
93                     break;

```

```

89         case 1:
90         case 0:
91             //left left
92             this->rotate_right(subtree_root_index);
93             break;
94         default:
95             throw std::domain_error(std::string("Unexpected balance factor
96                 with heavy left subtree: ")
97                                     +
98                                     std::to_string(left_child.balance_factor(this->nodes)));
99     }
100 } else if (std::abs(root_bal_fact > 2)) {
101     throw std::domain_error(std::string("Unexpected balance factor when
102         checking for heavy subtree: ")
103                             + std::to_string(root_bal_fact));
104 }
105 }
106 void do_validate_avl_balance(size_t subtree_root_index) const {
107     if (_DEBUG_) {
108         if (subtree_root_index == 0) return;
109         Node const& n = this->nodes[subtree_root_index];
110         if (abs(n.balance_factor(this->nodes)) > 1)
111             throw std::domain_error("Unexpected unbalanced tree while
112                 checking balance factor of all tree nodes");
113         do_validate_avl_balance(n.left_index);
114         do_validate_avl_balance(n.right_index);
115     }
116 }
117 void validate_avl_balance() {
118     if (_DEBUG_)
119         do_validate_avl_balance(this->root_index);
120 }
121 public:
122 AVL(size_t capacity): super(capacity) {}
123 /*
124     if there is space available, adds the specified key/value-pair to the
125     tree
126     and returns the number of nodes visited, V; otherwise returns -1 * V. If
127     an
128     item already exists in the tree with the same key, replace its value.
129 */
130 int insert(key_type const& key, value_type const& value) {
131     if (this->size() == this->capacity())
132         //no more space
133         return -1 * this->size();
134     bool found_key = false;
135     Key k(key);
136     Value v(value);
137     int nodes_visited = insert_at_leaf(0, this->root_index, k, v, found_key);
138     validate_avl_balance();
139     if (_DEBUG_)
140         this->nodes[this->root_index].validate_children_count_recursive(this->nodes);

```

```

135         return nodes_visited;
136     }
137     /*
138         if there is an item matching key, removes the key/value-pair from the
            tree, stores
139         it's value in value, and returns the number of probes required, V;
            otherwise returns -1 * V.
140     */
141     int remove(key_type const& key, value_type& value) {
142         if (this->is_empty())
143             return 0;
144         bool found_key = false;
145         Key k(key);
146         Value v(value);
147         int nodes_visited = do_remove(0, this->root_index, k, v, found_key);
148         validate_avl_balance();
149         if (_DEBUG_)
150             this->nodes[this->root_index].validate_children_count_recursive(this->nodes);
151         if (found_key)
152             value = v.raw_copy();
153         return found_key ? nodes_visited : -1 * nodes_visited;
154     }
155 };
156 }
157
158 #endif

```
